


After go func(): Goroutines Through a Beginner's Eye



Vaibhav Gupta


 [@97vaibhav](https://github.com/@97vaibhav)

- *Indian *
- *Backend engineer@Qenest Holdings*
- *A Three-year-old Gopher*
- *Second Time
Speaker@GoConference
Tokyo*

After go func(): Goroutines Through a Beginner's Eye

Outline

- *Motivation*
- *Goroutines & Go Scheduler Model*
- *Scheduler Internals (Fairness ,Preemption Work Stealing)*
- *Visualization*
- *Beginner Pitfalls & My Learnings*
- *Conclusion*
- *References*
- *Q/A*



Motivation

- *One keyword, massive power (go func())*
1つのキーワードで、ものすごい力 (go func())
- *From “it works” to “I understand why”*
「なぜか動く」から「理解して動かす」へ
- *Today’s goal: a clear mental model*
今日の目標: 頭の中にクリアなメンタルモデルを描く





Goroutine は何?



Goroutines are lightweight threads which are managed by Go runtime
ゴルーチンとは、Go ランタイムに管理される軽量スレッド。


Goroutines





<i>Goroutines</i>	<i>OS Threads</i>
<ol style="list-style-type: none">1. <i>language-level, managed by Go runtime</i>2. <i>very cheap to create 2kb</i>3. <i>growable segmented stacks</i>4. <i>millions possible; parallelism limited by GOMAXPROCS (P count)</i>	<ol style="list-style-type: none">1. <i>kernel-level, managed by the OS scheduler</i>2. <i>expensive to create</i>3. <i>fixed-size stacks</i>4. <i>hundreds–thousands practical; parallelism capped by CPU cores</i>

Goroutines を見ましょう ...



```
package main

import (
    "fmt"
    "sync"
)

func main() {
    var wg sync.WaitGroup
    wg.Add(11)
    for i := 0; i <= 10; i++ {
        go func(i int) {
            defer wg.Done()
            fmt.Printf("Printing value of i in goroutine -- %d\n", i)
        }(i)
    }
    wg.Wait()
    fmt.Println("Hello, Welcome to Goroutines")
}
```

go run demo1.go

1st Run

```
package main

import (
    "fmt"
    "sync"
)

func main() {
    var wg sync.WaitGroup
    wg.Add(11)
    for i := 0; i <= 10; i++ {
        go func(i int) {
            defer wg.Done()
            fmt.Printf("Printing value of i in goroutine -- %d\n", i)
        }(i)
    }
    wg.Wait()
    fmt.Println("Hello, Welcome to Goroutines")
}
```

go run demo1.go

```
~/golang_projects/github.com/97vaibhav/go-conference-2025/demo1
> go run demo1.go
Printing value of i in goroutine - 10
Printing value of i in goroutine - 4
Printing value of i in goroutine - 7
Printing value of i in goroutine - 1
Printing value of i in goroutine - 2
Printing value of i in goroutine - 3
Printing value of i in goroutine - 8
Printing value of i in goroutine - 9
Printing value of i in goroutine - 5
Printing value of i in goroutine - 6
Printing value of i in goroutine - 0
Hello, Welcome to Goroutines
```

2nd Run

go run demo1.go

```
package main

import (
    "fmt"
    "sync"
)

func main() {
    var wg sync.WaitGroup
    wg.Add(11)
    for i := 0; i <= 10; i++ {
        go func(i int) {
            defer wg.Done()
            fmt.Printf("Printing value of i in goroutine -- %d\n", i)
        }(i)
    }
    wg.Wait()
    fmt.Println("Hello, Welcome to Goroutines")
}
```

```
~/golang_projects/github.com/97vaibhav/go-conference-2025/demo1
> go run demo1.go
Printing value of i in goroutine - 1
Printing value of i in goroutine - 10
Printing value of i in goroutine - 0
Printing value of i in goroutine - 3
Printing value of i in goroutine - 4
Printing value of i in goroutine - 5
Printing value of i in goroutine - 6
Printing value of i in goroutine - 7
Printing value of i in goroutine - 8
Printing value of i in goroutine - 9
Printing value of i in goroutine - 2
Hello, Welcome to Goroutines
```

3rd Run

go run demo1.go

```
package main

import (
    "fmt"
    "sync"
)

func main() {
    var wg sync.WaitGroup
    wg.Add(11)
    for i := 0; i <= 10; i++ {
        go func(i int) {
            defer wg.Done()
            fmt.Printf("Printing value of i in goroutine -- %d\n", i)
        }(i)
    }
    wg.Wait()
    fmt.Println("Hello, Welcome to Goroutines")
}
```

```
~/golang_projects/github.com/97vaibhav/go-conference-2025/demo1
> go run demo1.go
Printing value of i in goroutine - 10
Printing value of i in goroutine - 3
Printing value of i in goroutine - 0
Printing value of i in goroutine - 1
Printing value of i in goroutine - 2
Printing value of i in goroutine - 4
Printing value of i in goroutine - 5
Printing value of i in goroutine - 8
Printing value of i in goroutine - 7
Printing value of i in goroutine - 6
Printing value of i in goroutine - 9
Hello, Welcome to Goroutines
```

- *How did 11 goroutines run concurrently? Magic?*
11個のgoroutineがどうやって並行して実行されたのでしょうか？魔法でしょうか？
- *In What Order 11 goroutines ran?*
11 個の goroutine はどのような順序で実行されましたか？


“If I Tell you in what order
goroutines run, it won’t happen” -
Dr Gopher Strange.

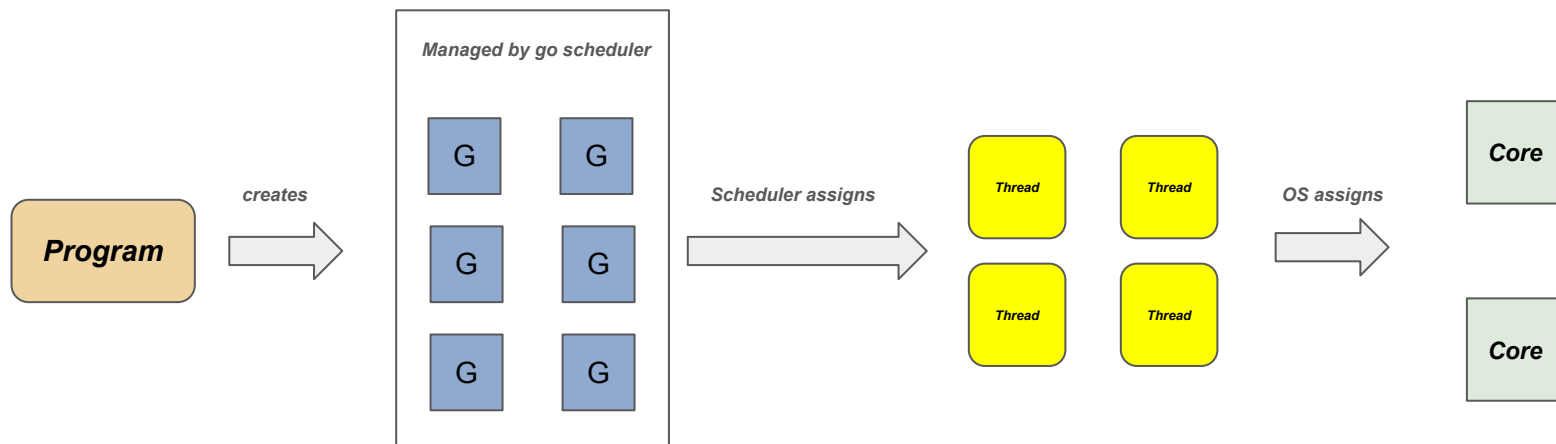


Go Scheduler Model




- *We need some way to map goroutines onto os threads- **User Space Scheduling***
ゴルーチンをOSスレッドにマッピングする方法が必要です - ユーザー空間スケジューリング






- *We need some way to map goroutines onto os threads- **User Space Scheduling***
ゴルーチンをOSスレッドにマッピングする方法が必要です - ユーザー空間スケジューリング



M:N Scheduling

- The no of G can be greater than number of M
Goroutine(G) の数は
Thread(M) より多くてよい
- Go scheduler multiplexes G onto available M
Go スケジューラが賢くMにGを
割り当ててくれる



Go scheduler Internals

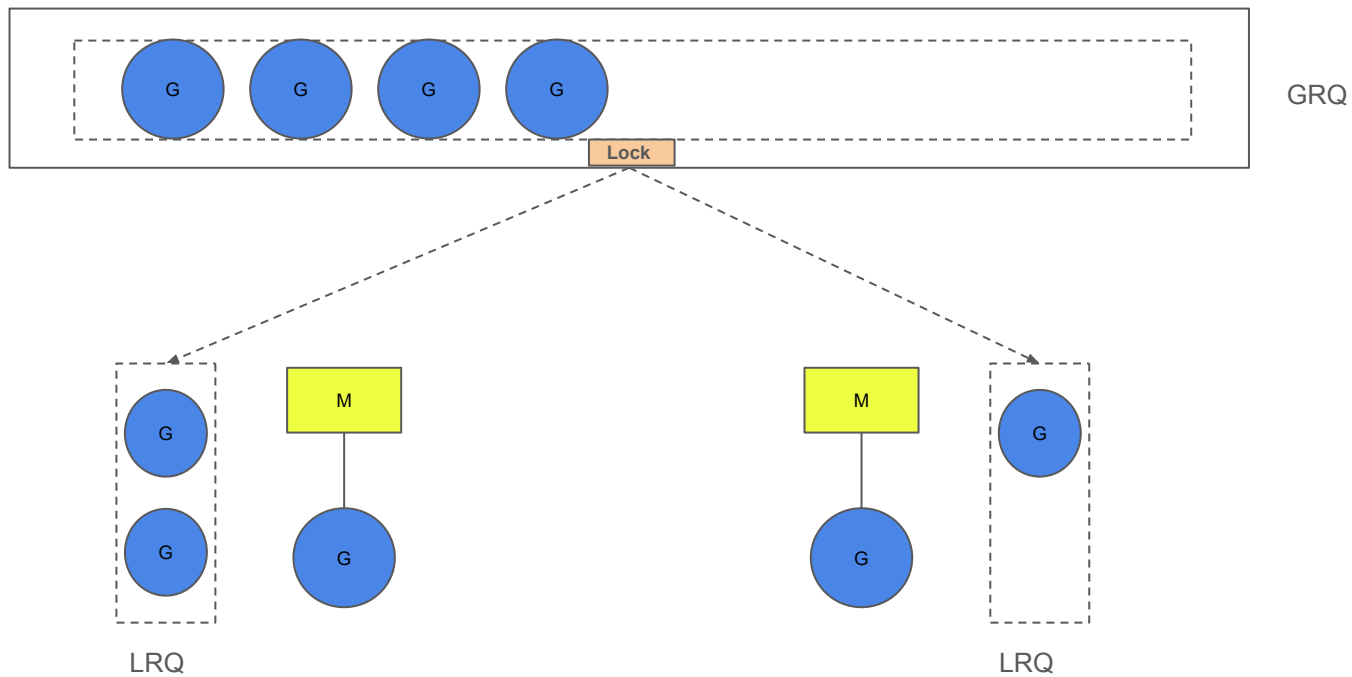
*How do we keep track of goroutine that are yet to be run
or are running ?*

まだ実行されていないゴルーチンを、どうやって管理しますか？



*Go*のスケジューラには、2つの実行キューがあります


- *Global Run Queue (GRQ)*
- *Local Run Queue (LRQ)*




GRQ = Global Run Queue
LRQ = Local Run Queue

But Wait there's a Problem !!


Long running tasks or System Calls ?? 😓



Processor



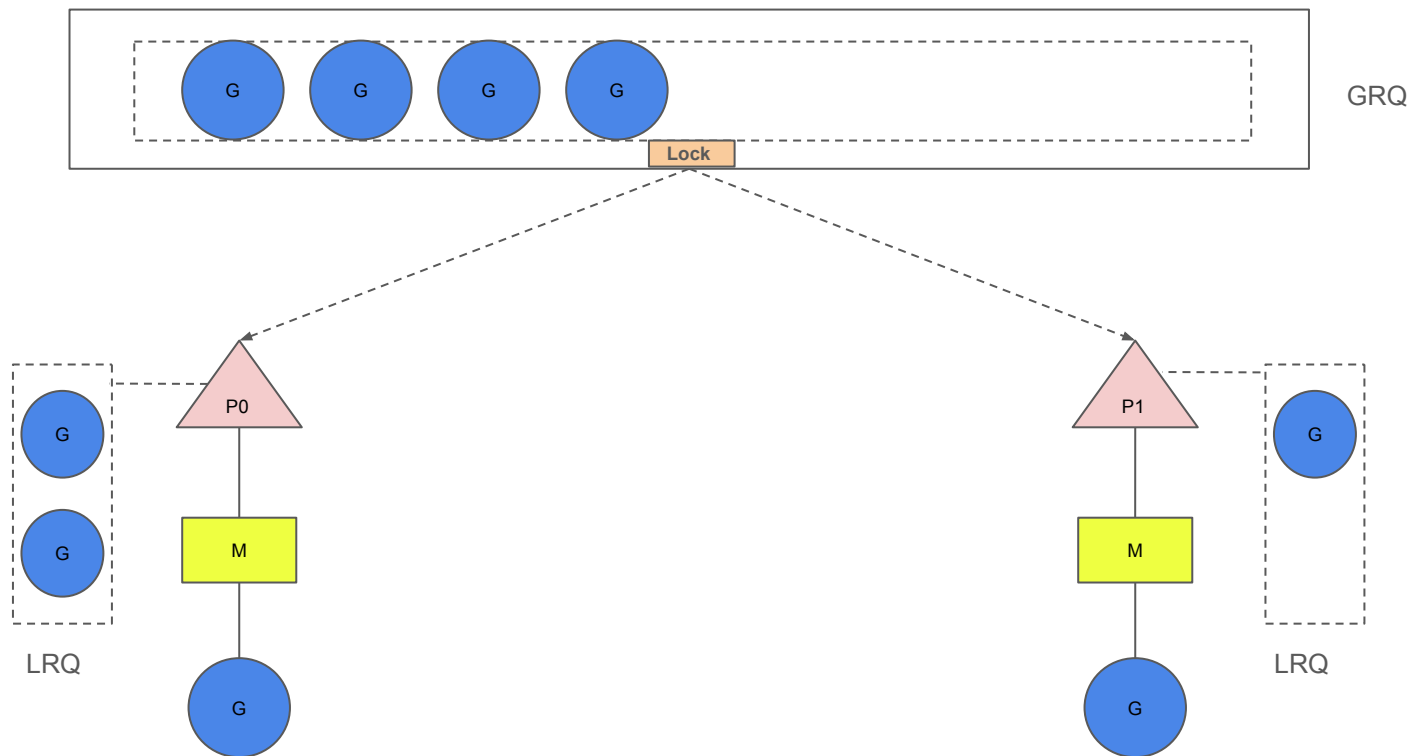
Processor



System Thread



Goroutine




GRQ = Global Run Queue
LRQ = Local Run Queue

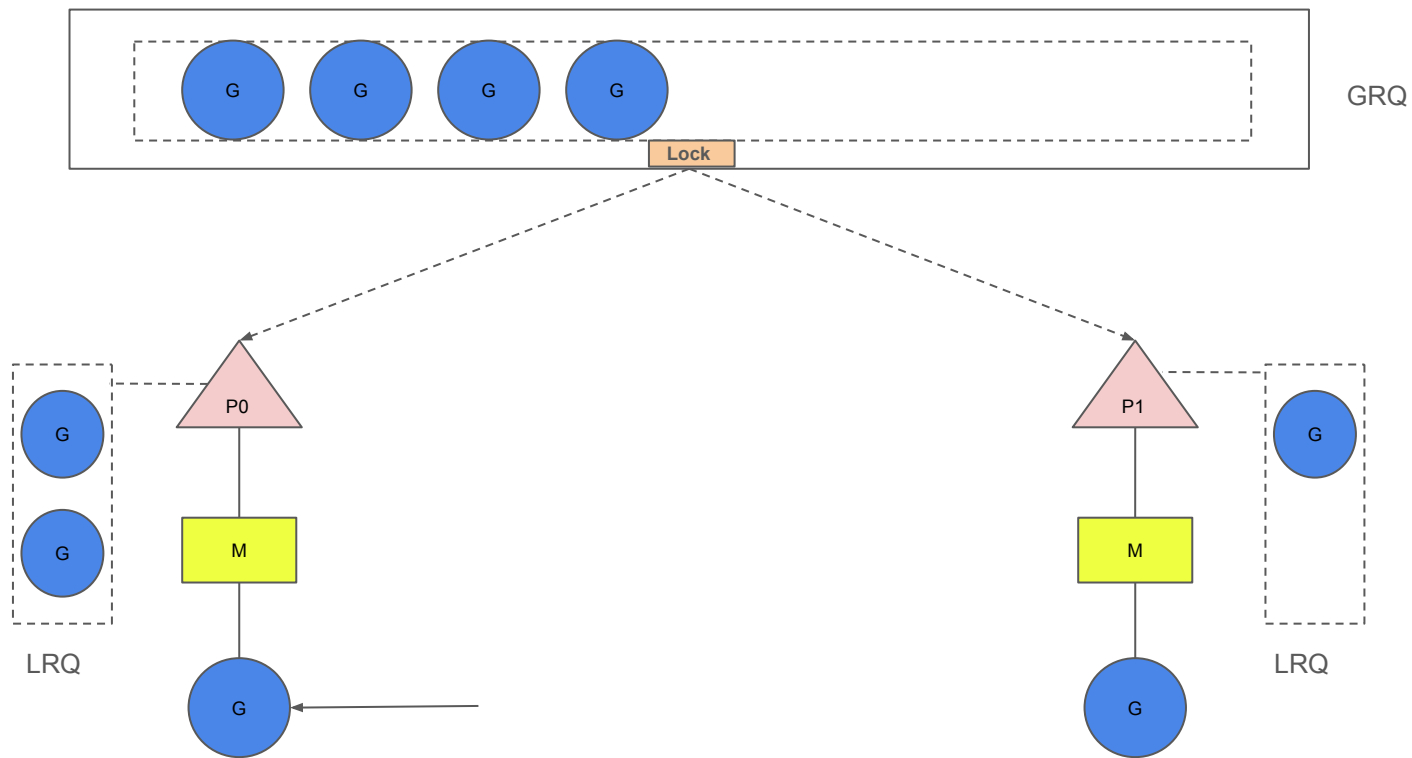
✨ プロセッサの数は最大数GOMAXPROCSです

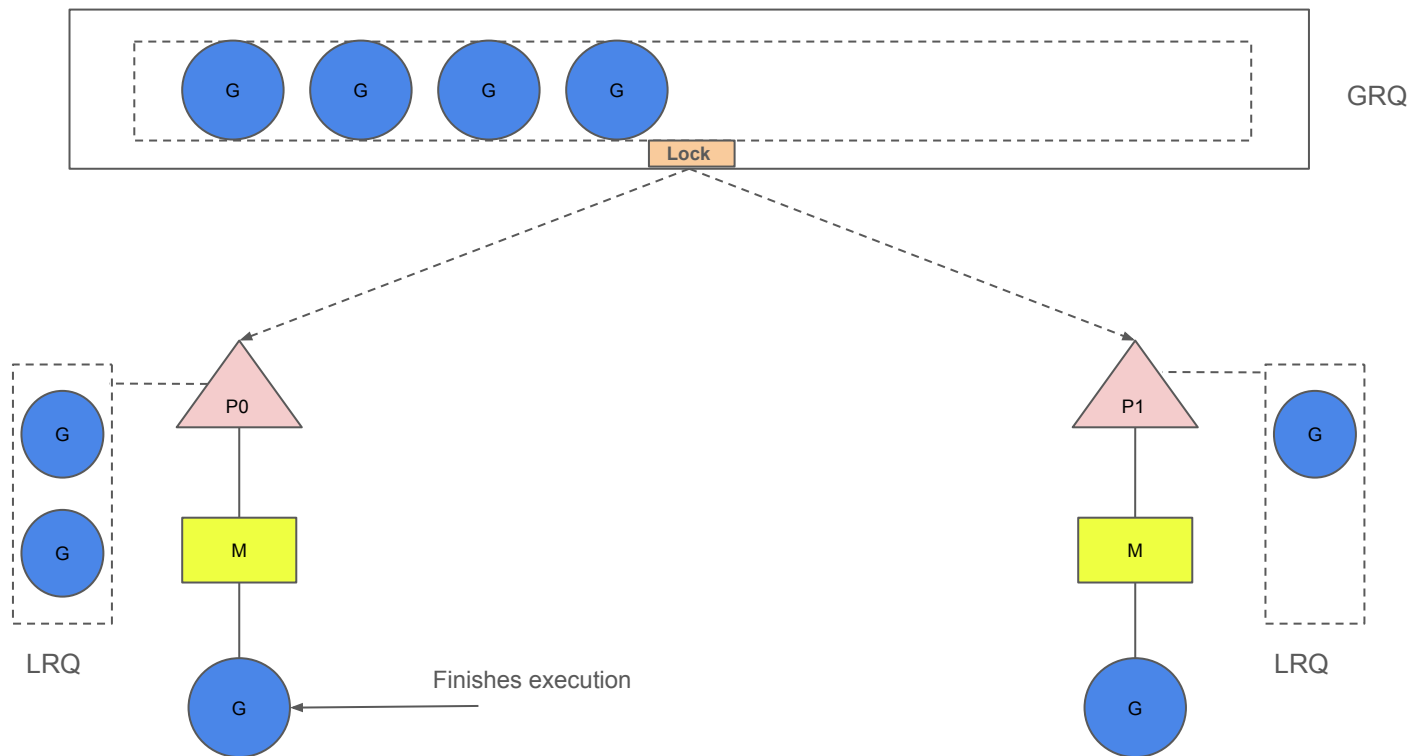
✨ *Now we have enough explanation*

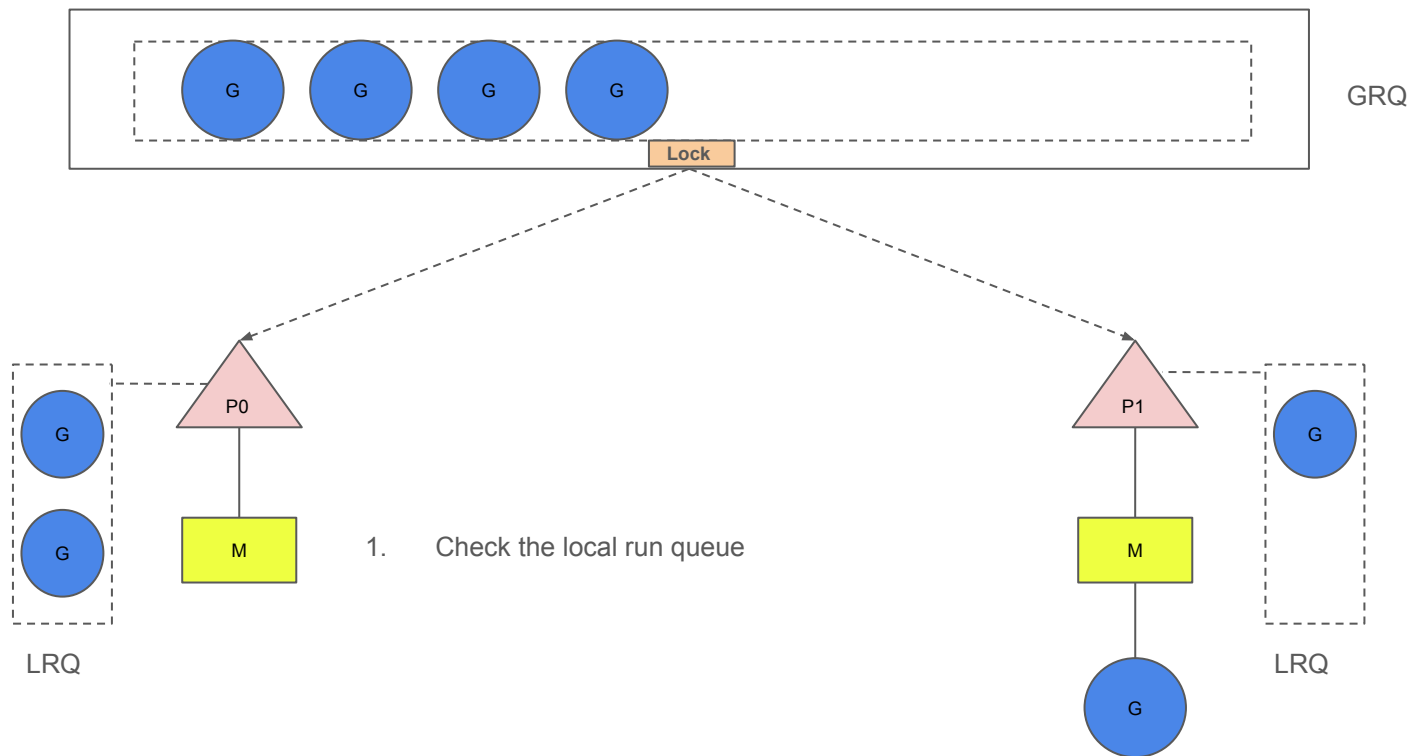
“How do we choose which go routine to run“

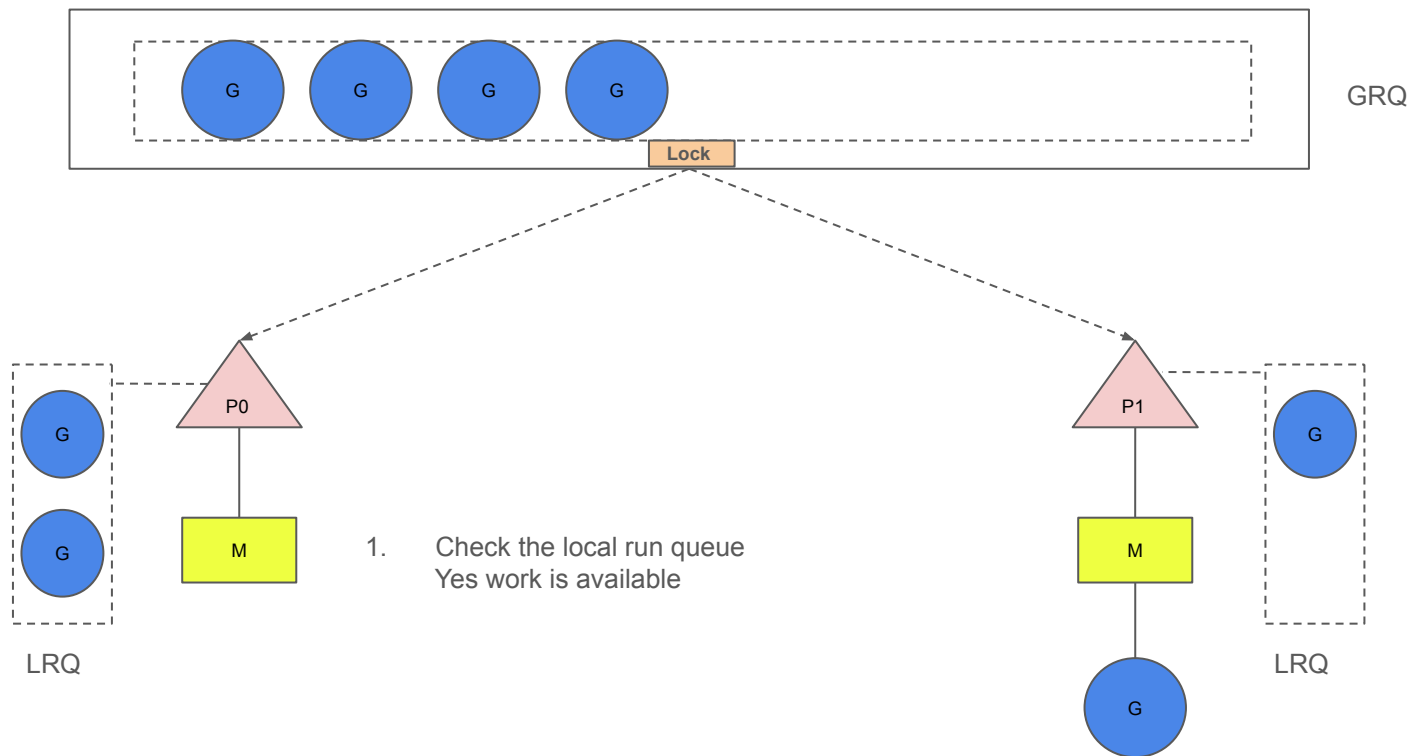


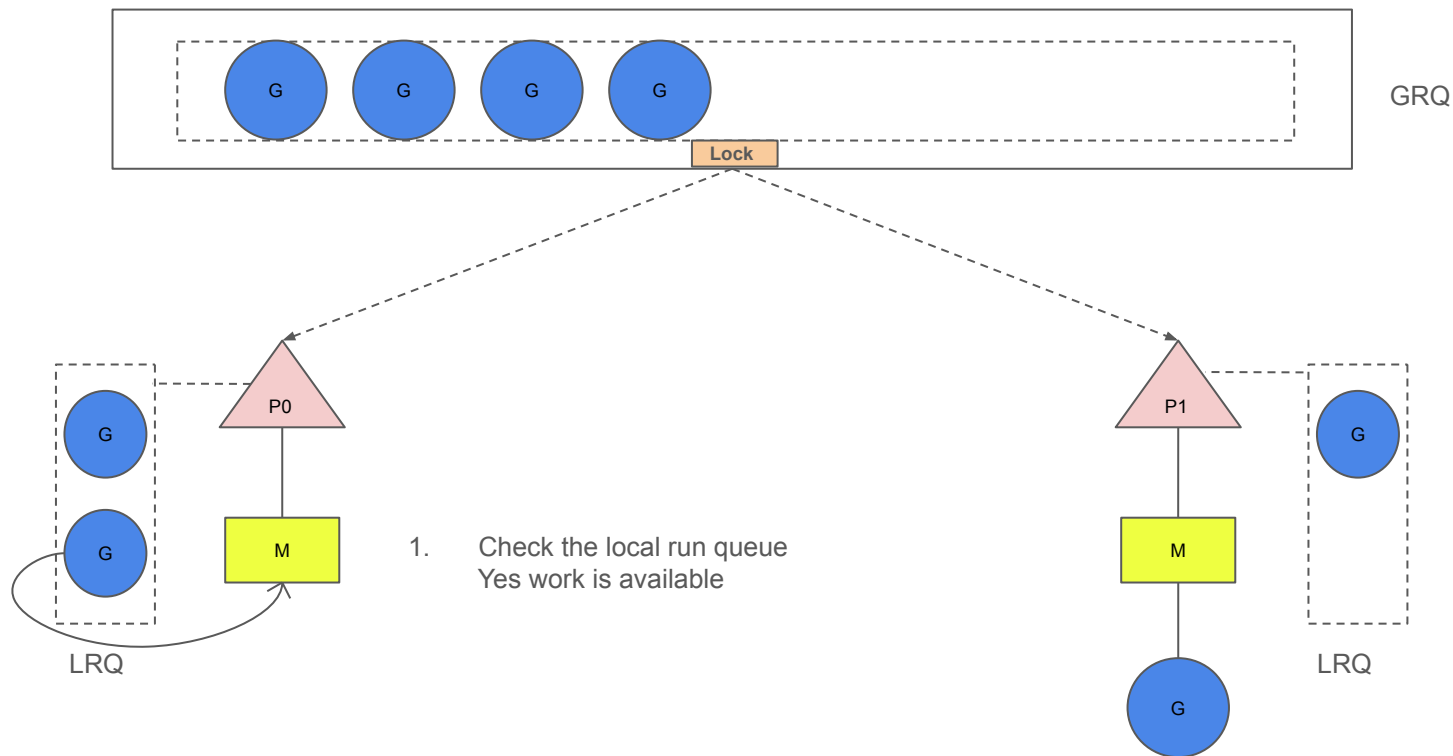
GRQ = Global Run Queue
LRQ = Local Run Queue

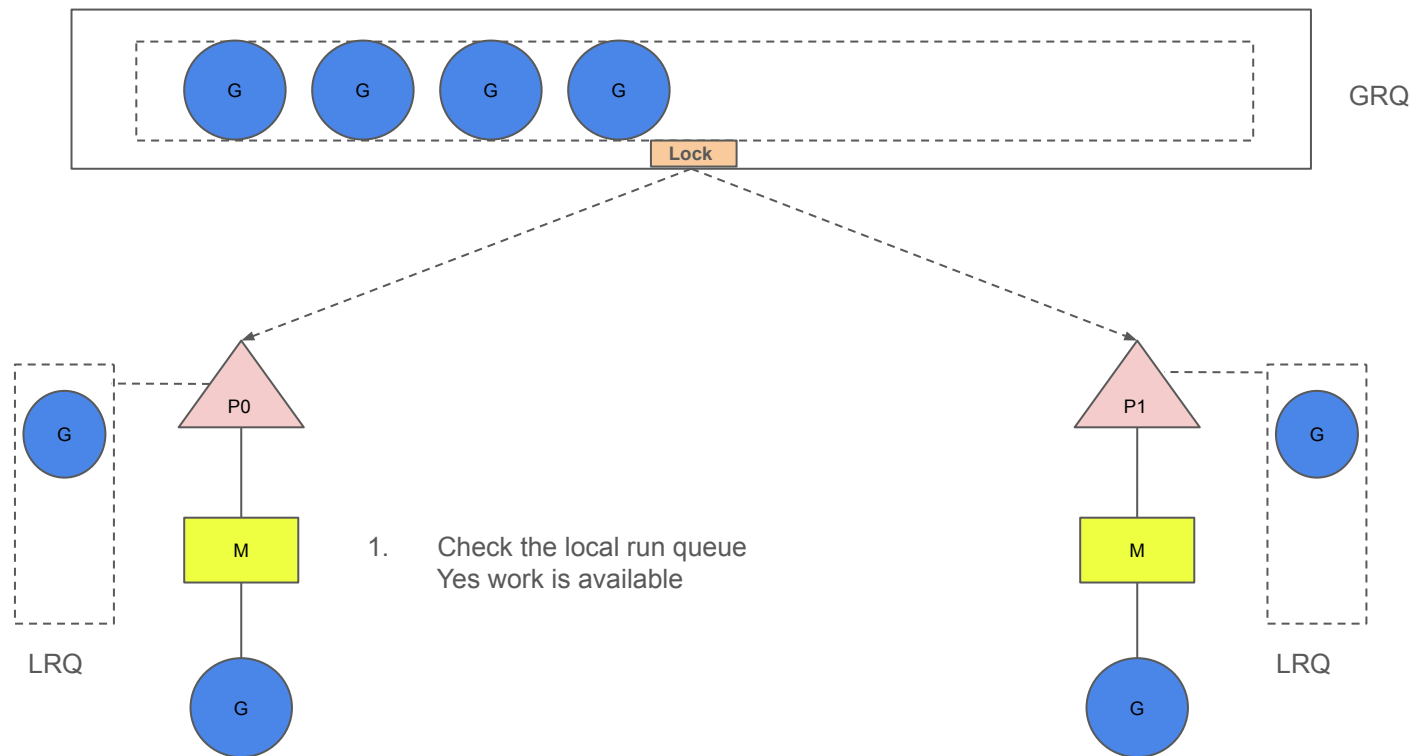




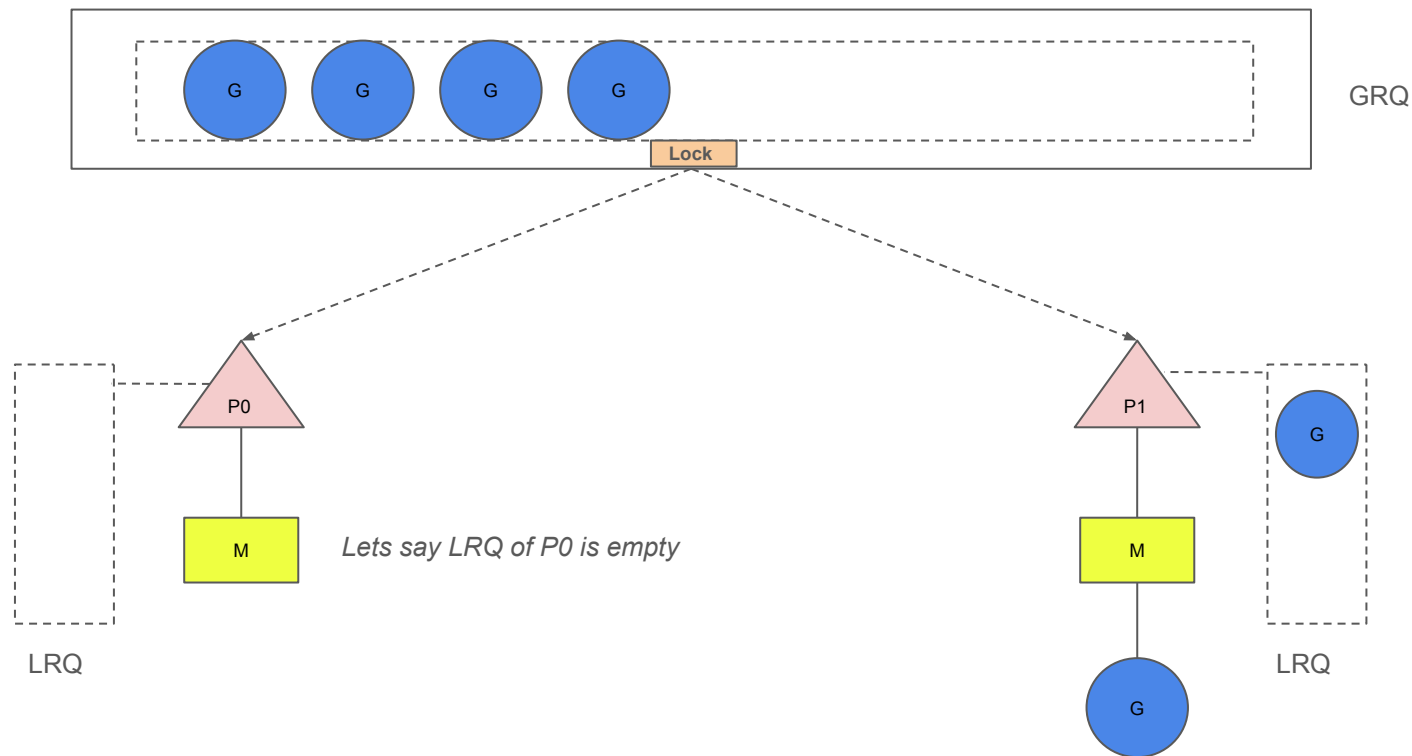


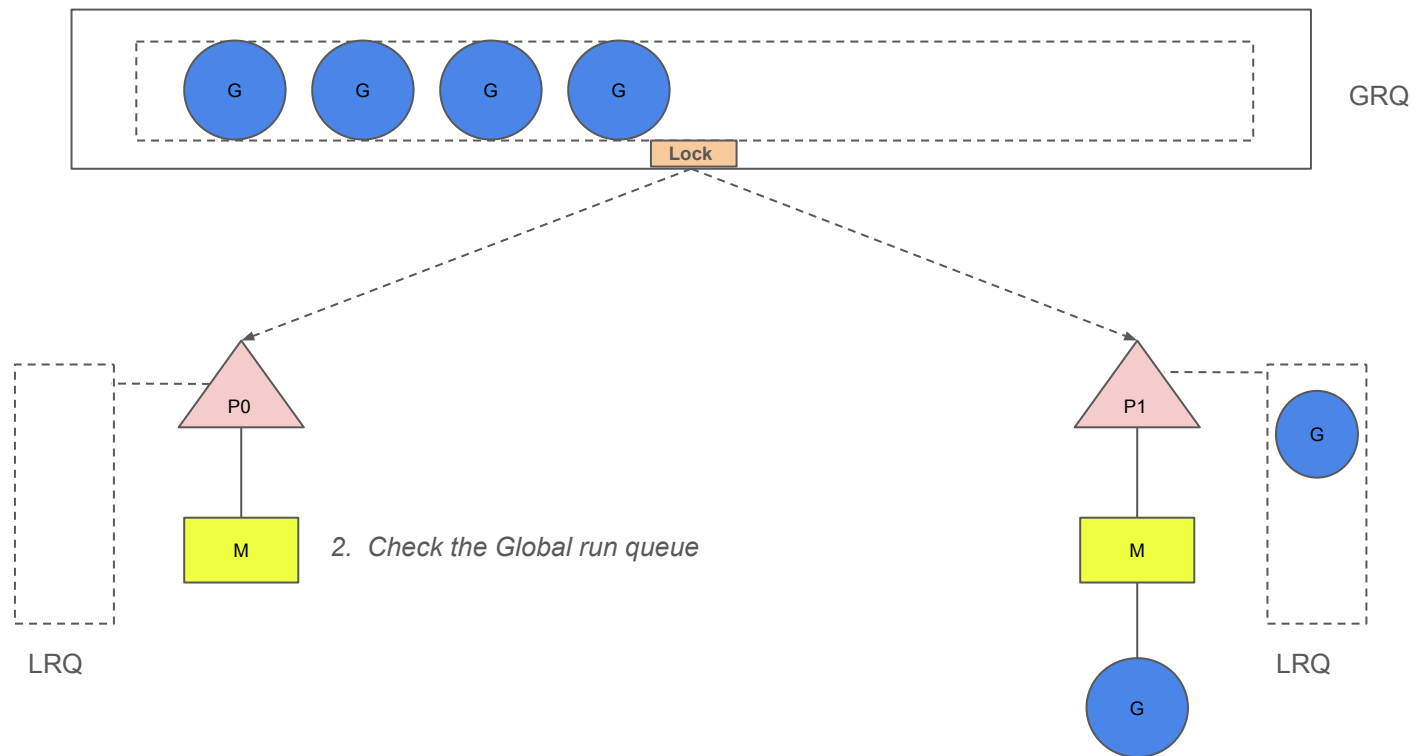


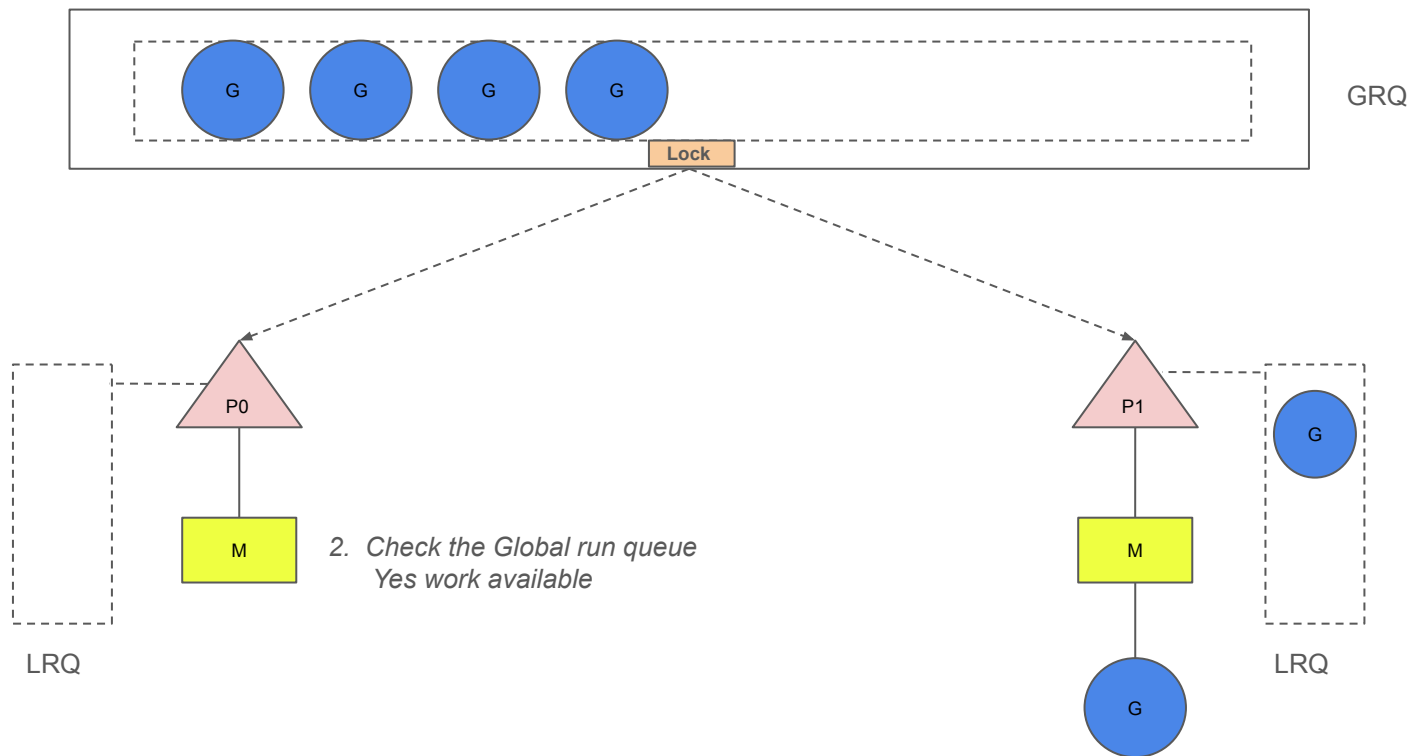


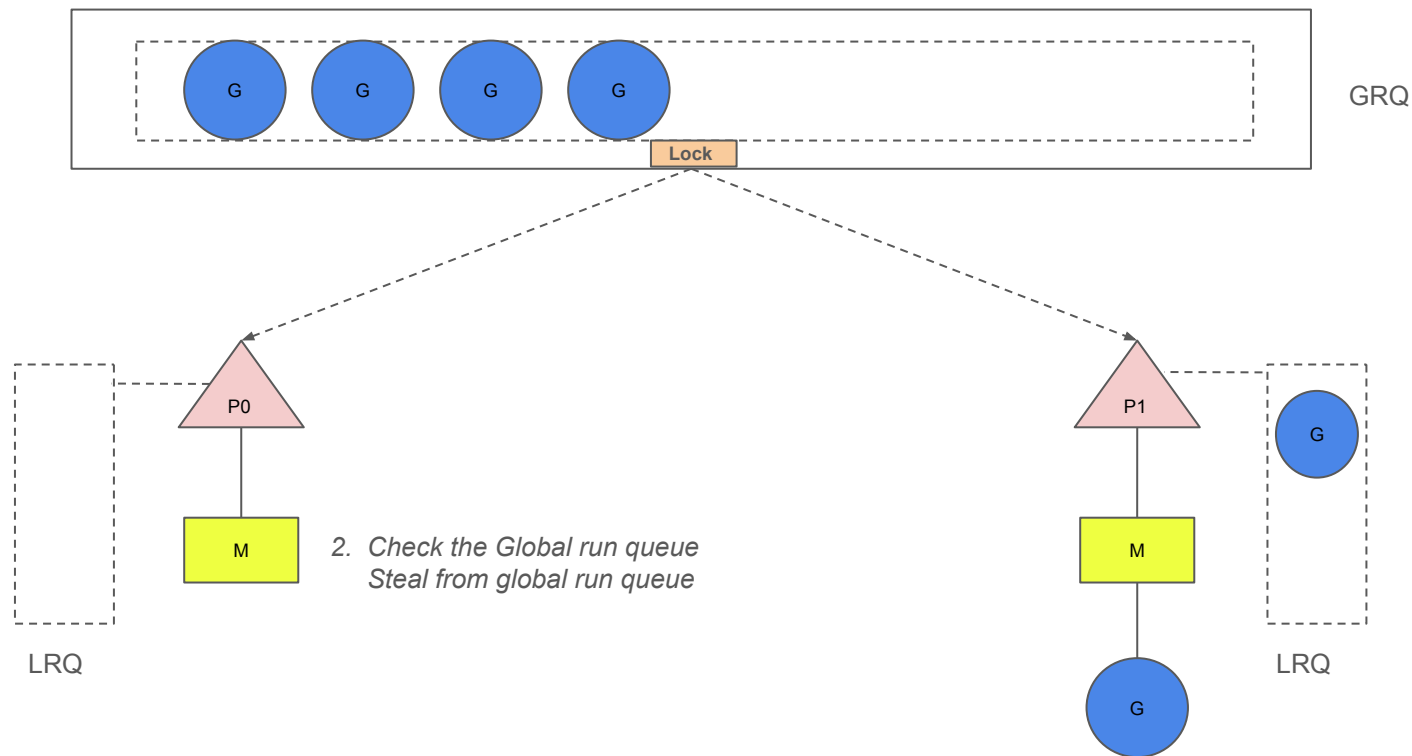


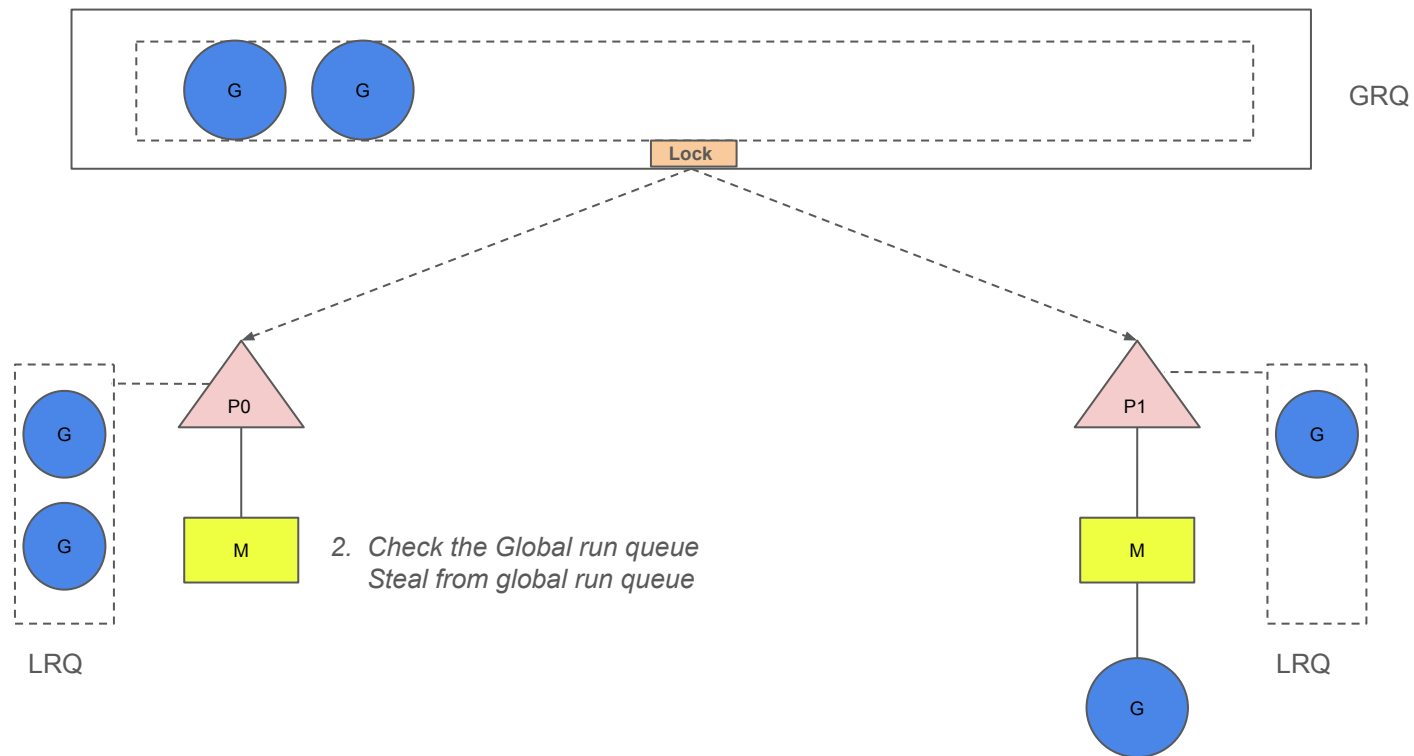
Local Run Queue は空ですか? 😅

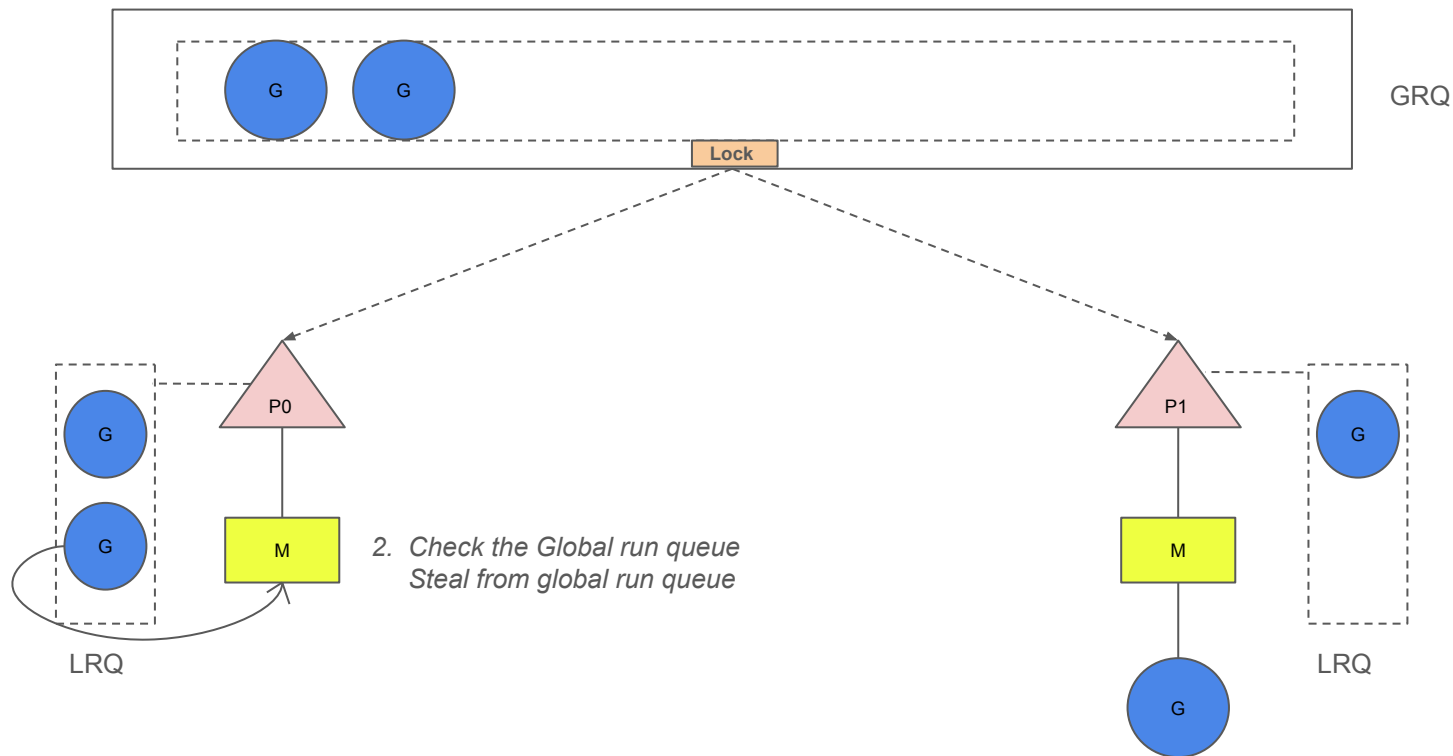


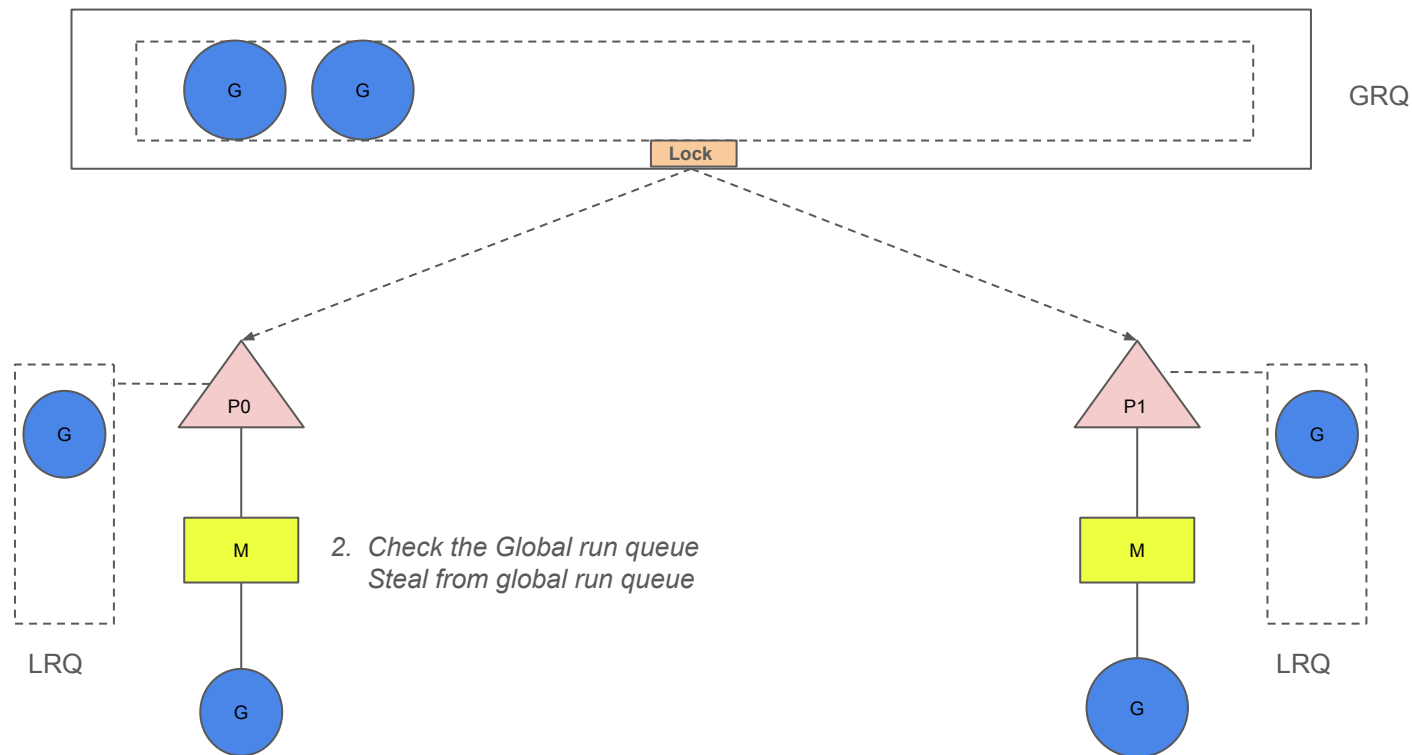




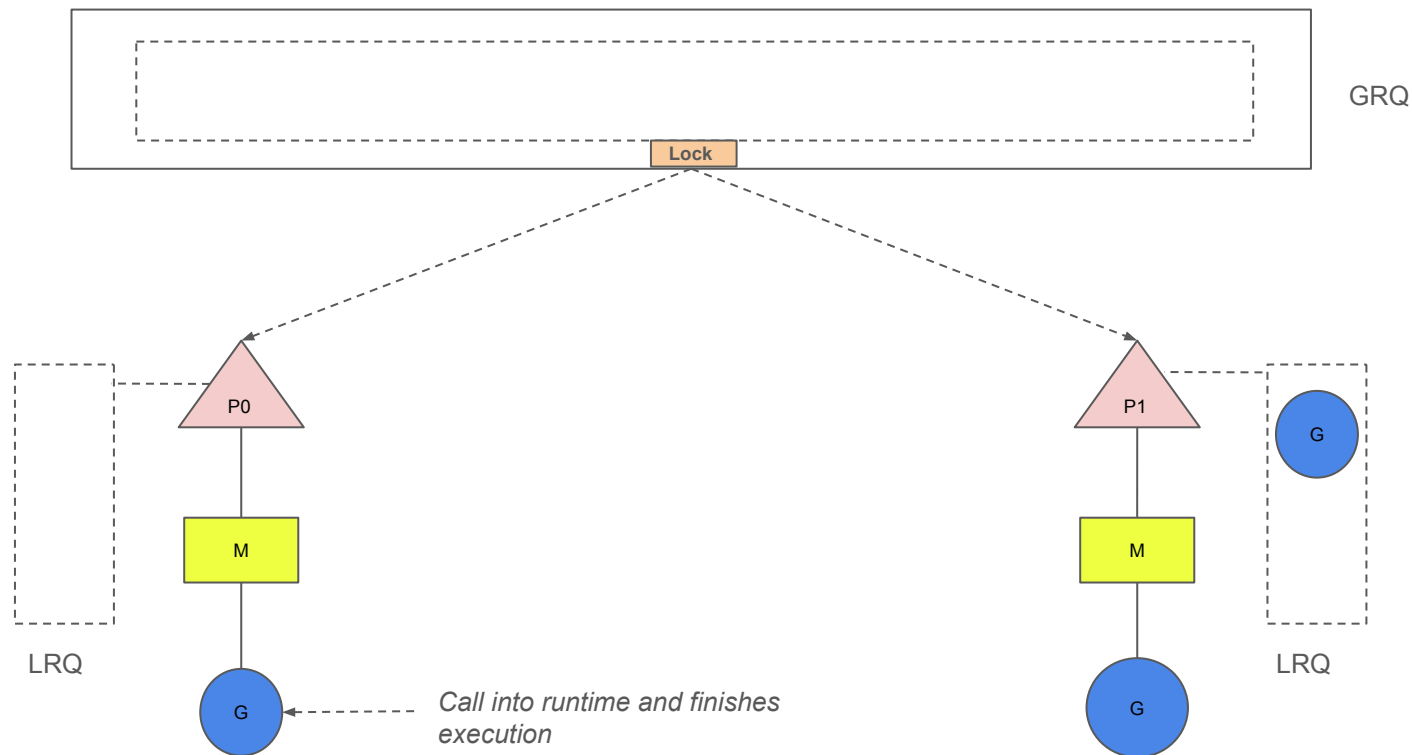


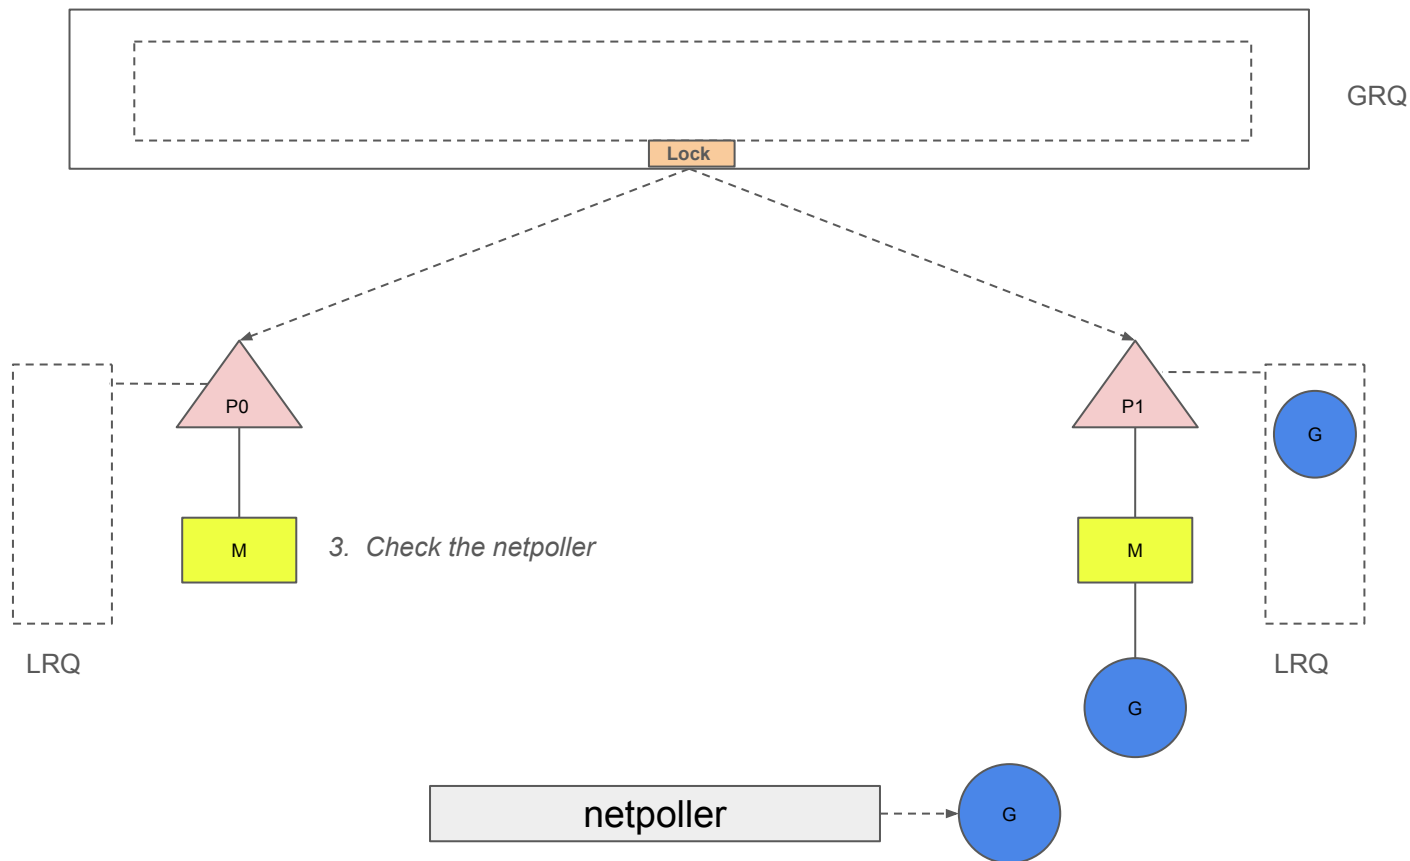


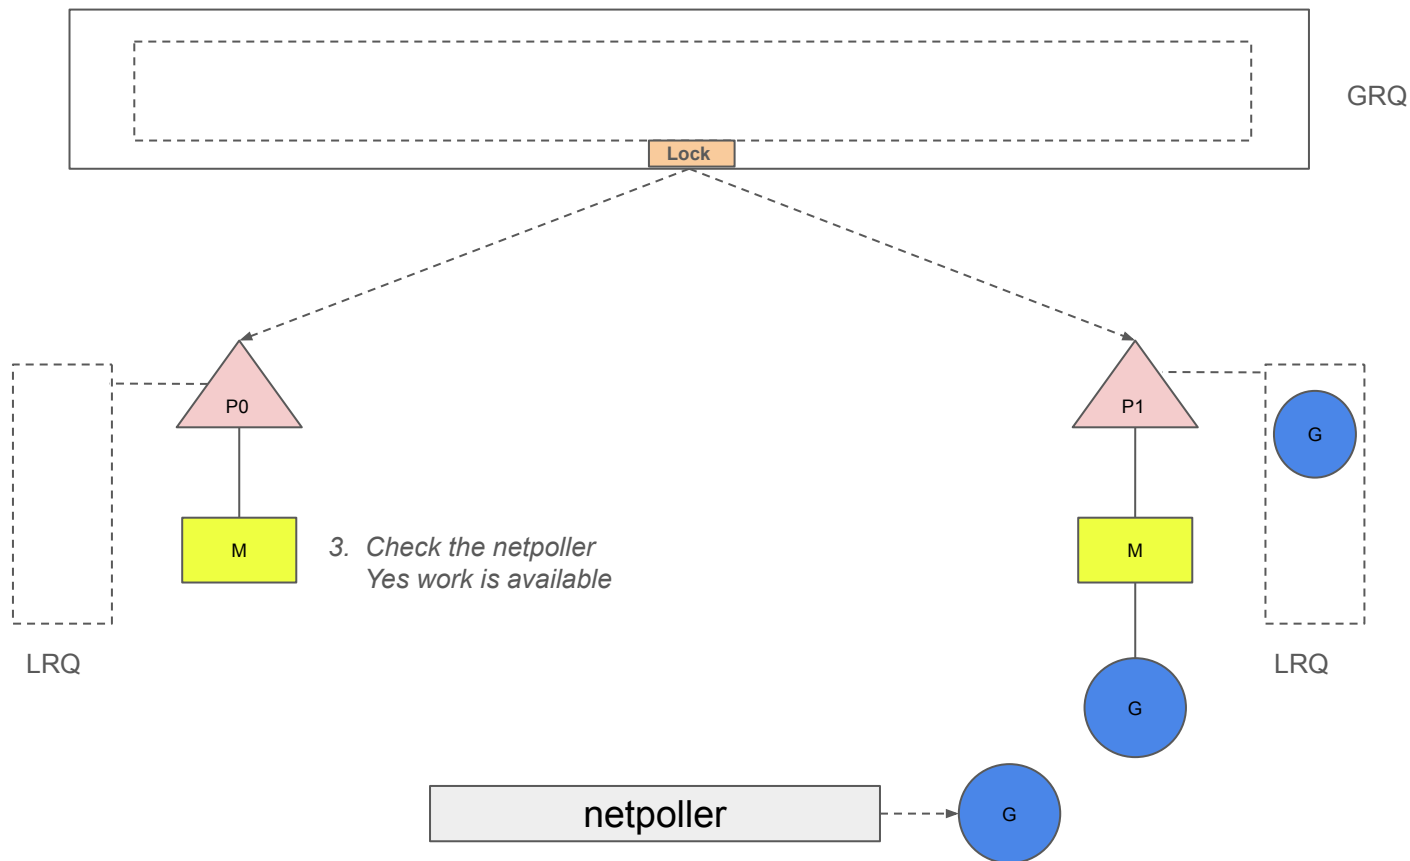


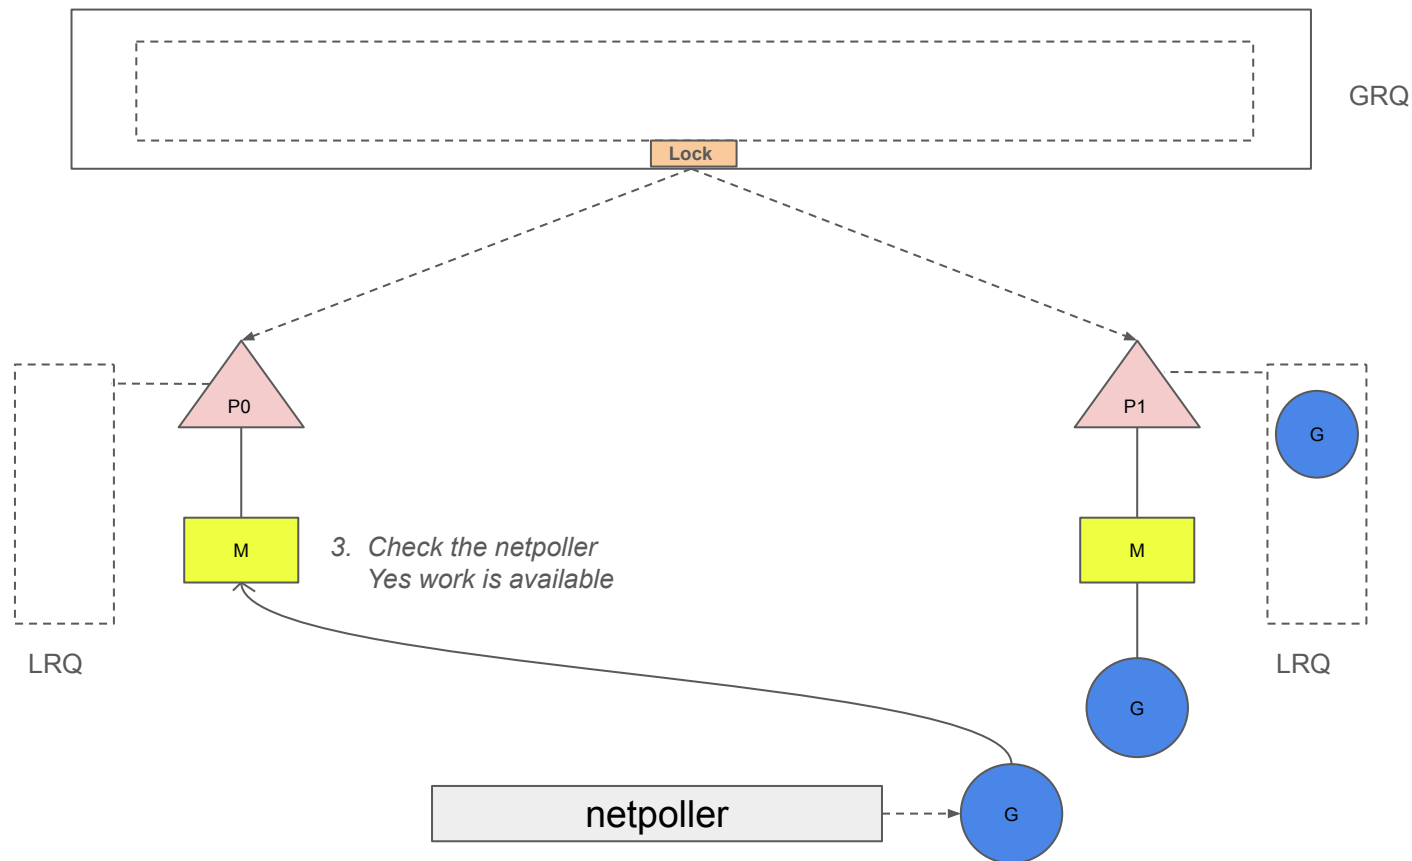


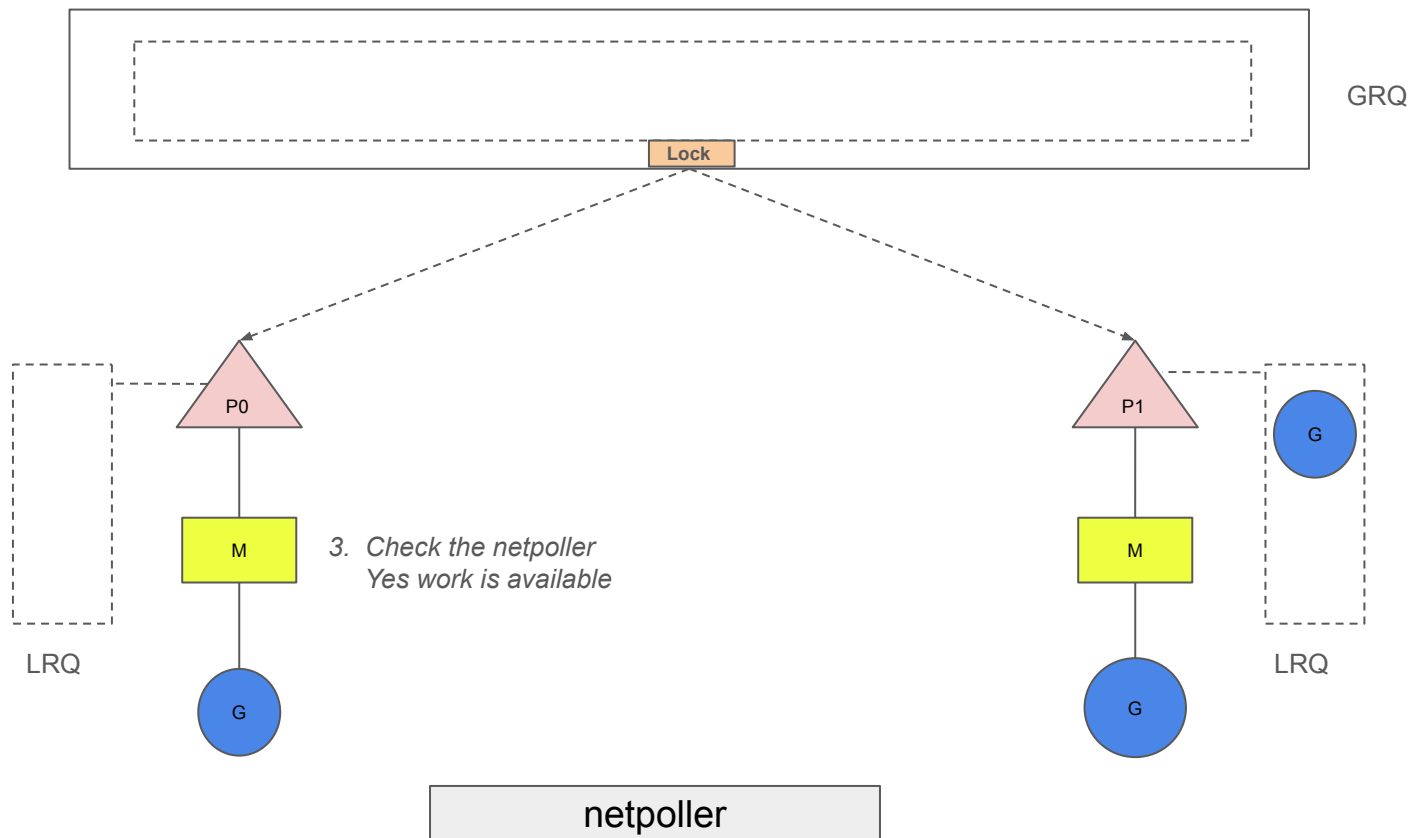
Global Run Queue 実行キューも空ですか？😓








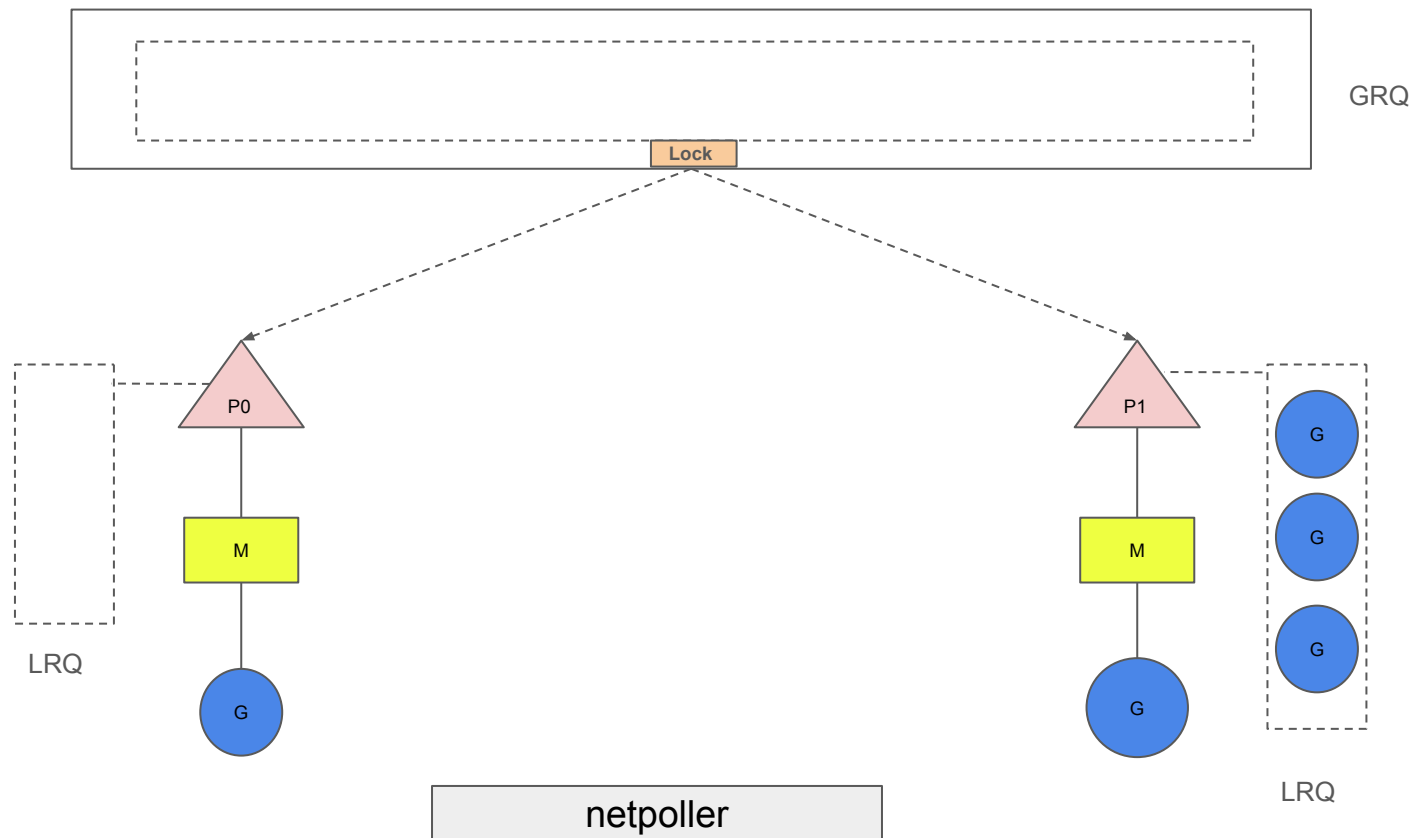


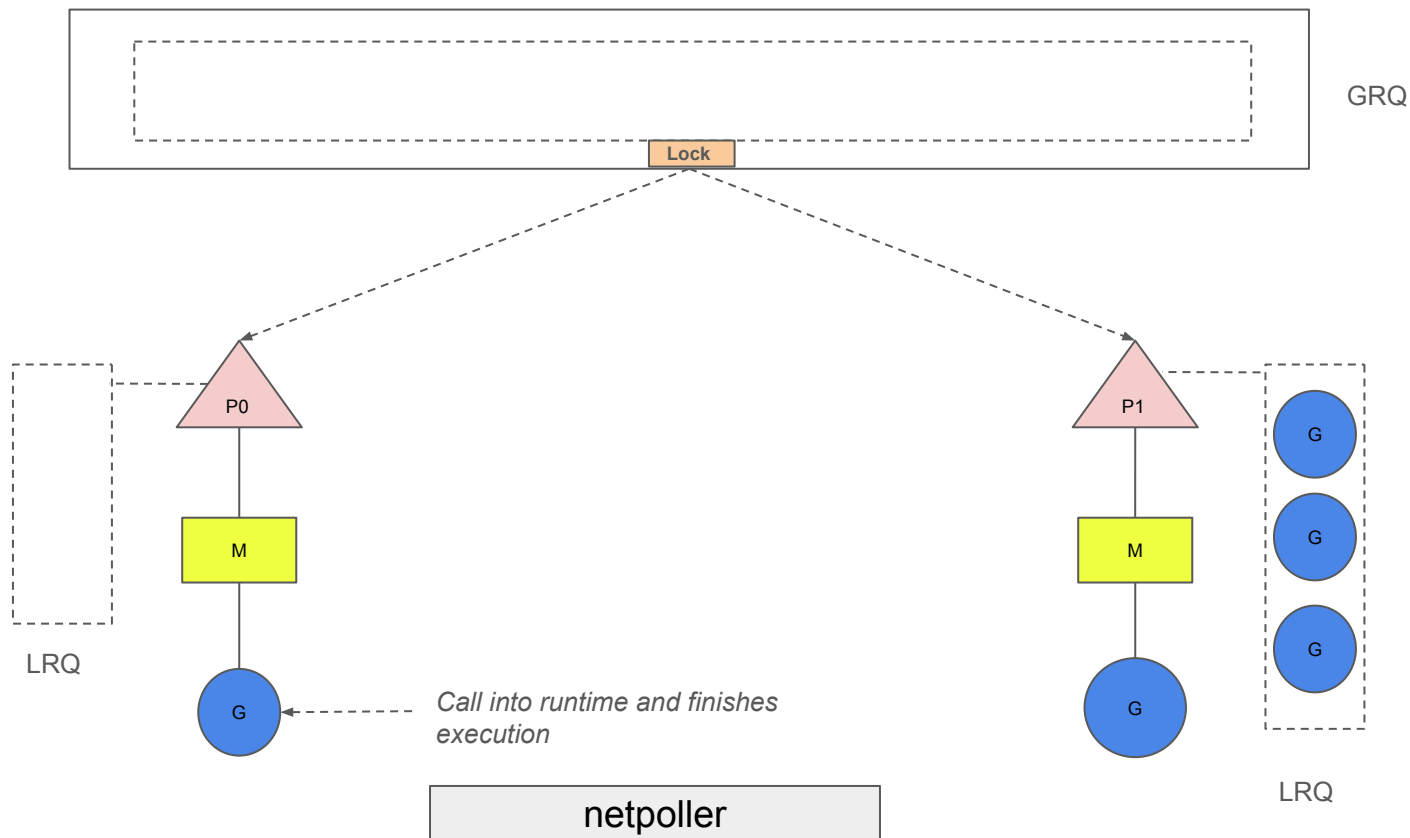


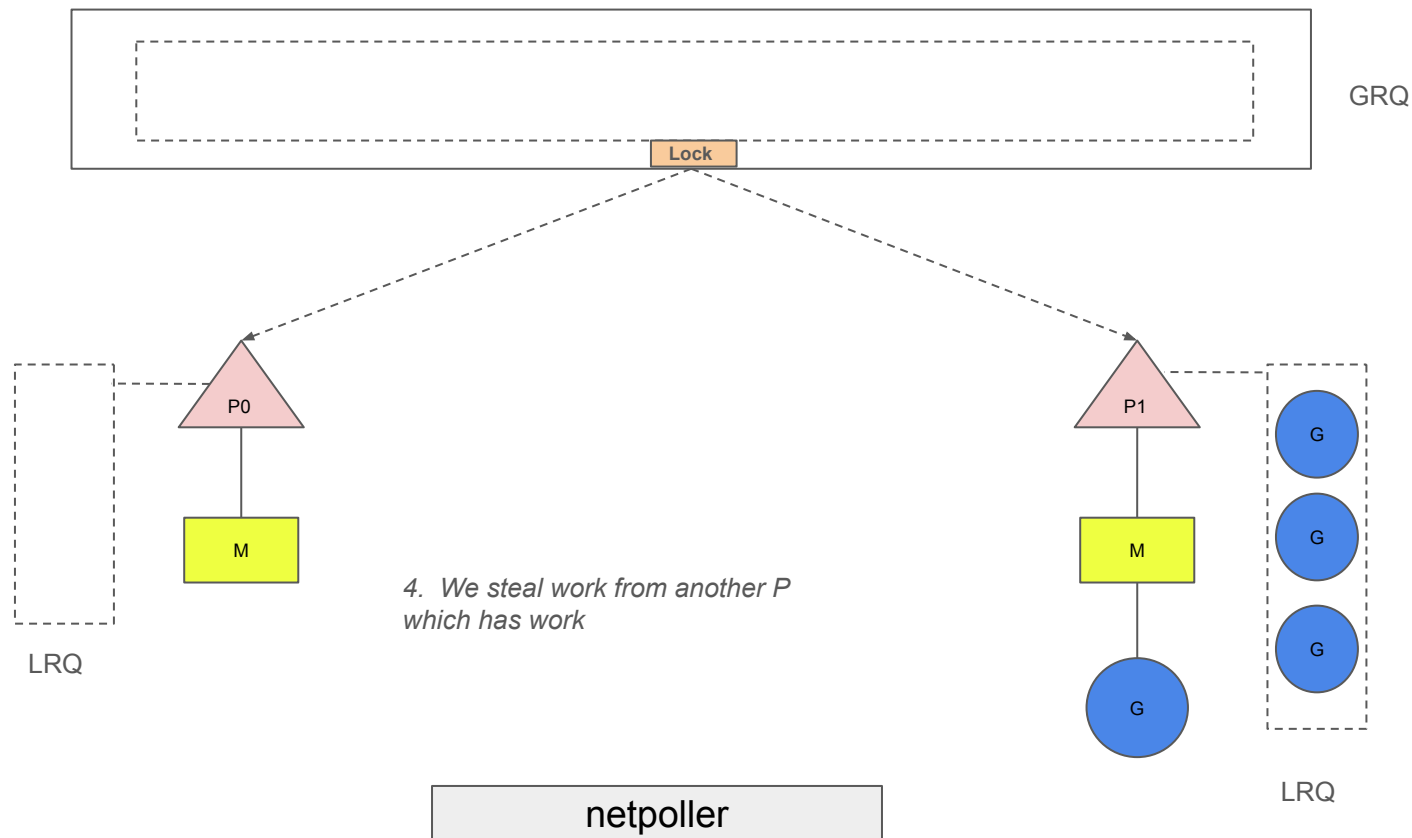
Netpoller も空の場合はどうなるでしょうか? 🤔

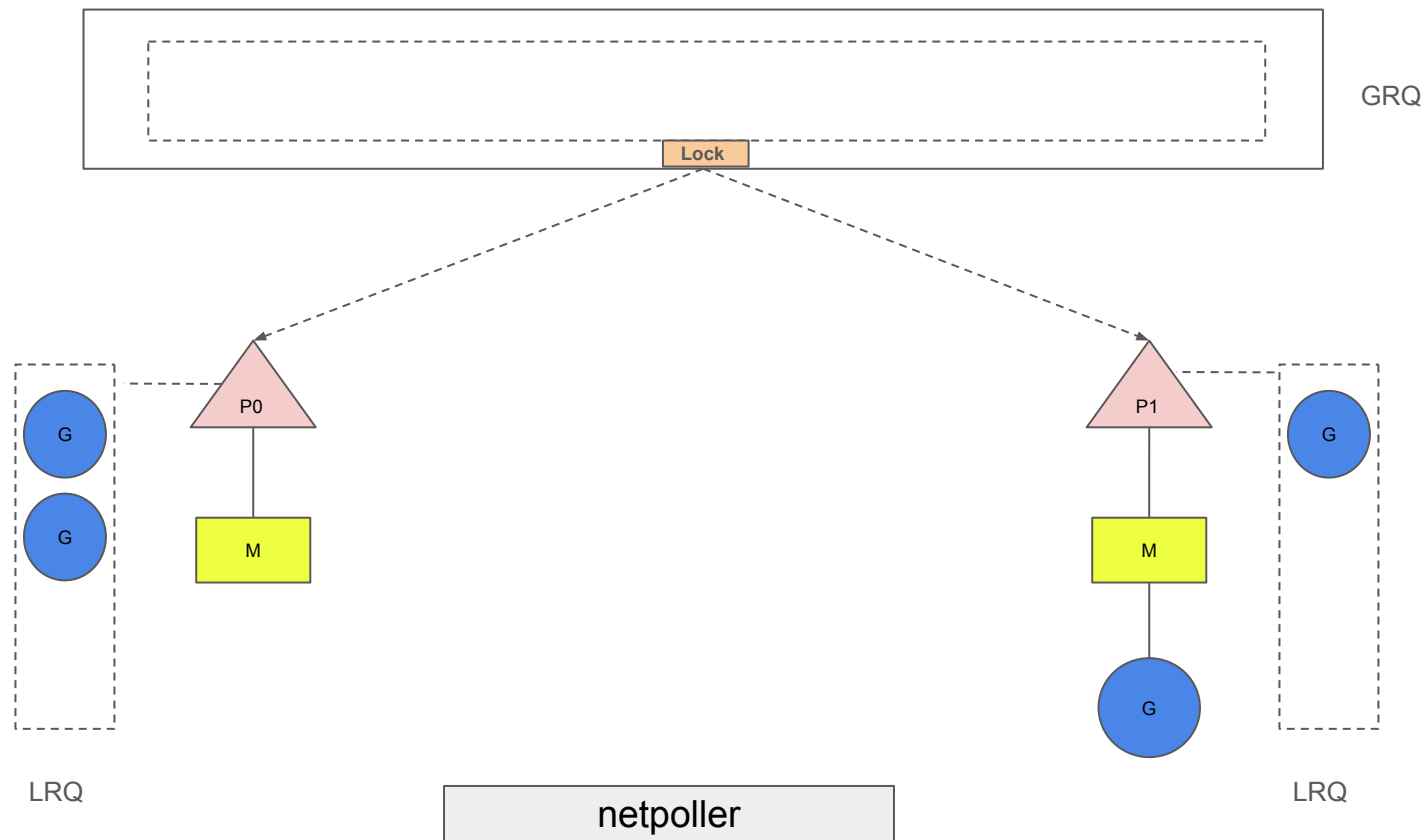
Work Stealing

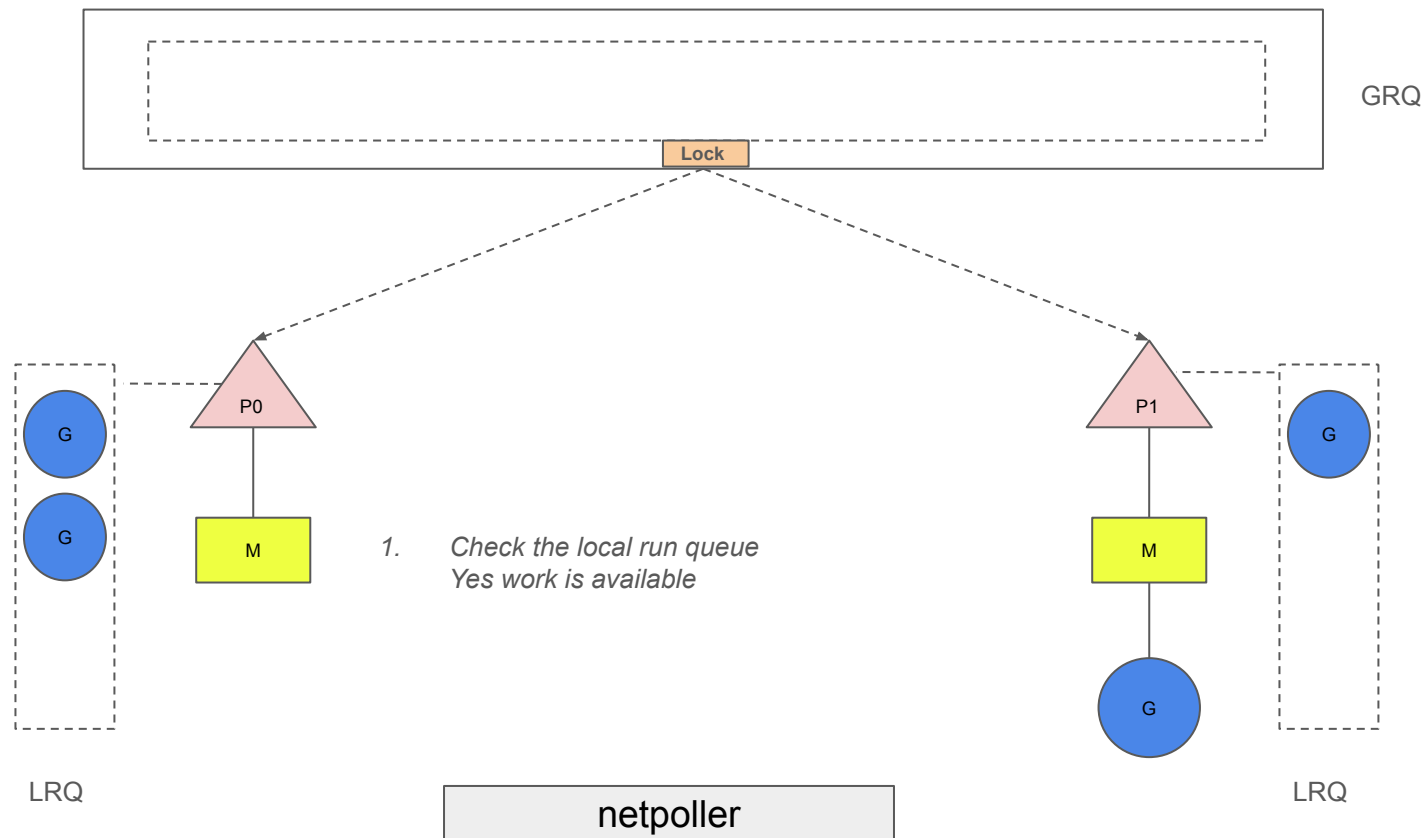


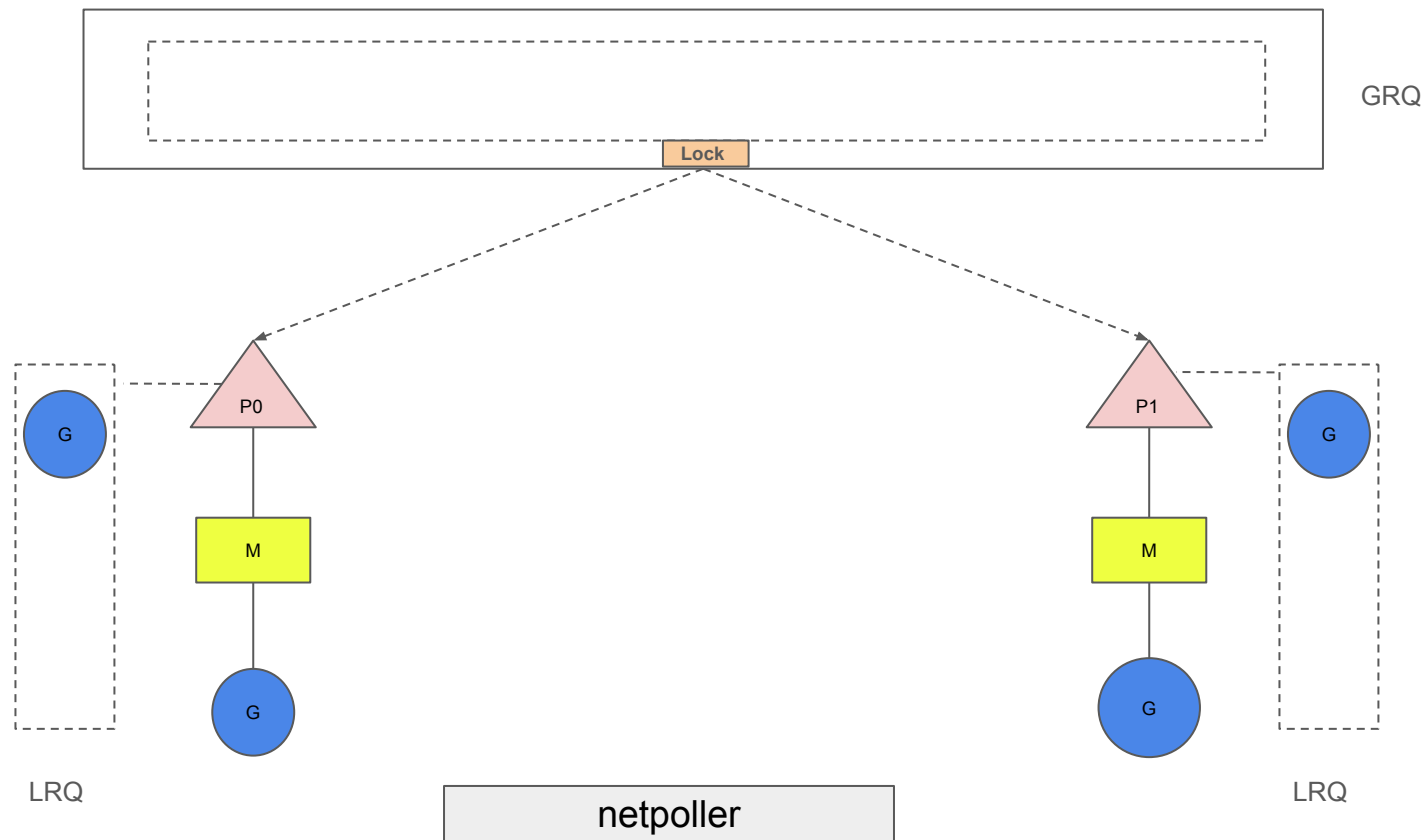














*Reference to the order of work
stealing execution -
[runtime/proc.go](https://runtime.proc.go)*

What about long running task

Preemption









Upto Go 1.10

- *Go has used cooperative preemption with safe-points only at function calls.*
- *From execution point of view you can give goroutine processor time (execution time) only on specific events (safe-points) which are function calls.*




What's Now ...



Non-Cooperative Preemption

- *Go introduced non-cooperative preemption because of the problems mentioned above.*
Go は、先ほどの問題を解決するために「非協調型プリエンプション」を導入しました。
- *In non-cooperative preemption, the Go runtime can forcibly pause a running goroutine even if it doesn't explicitly yield control. This preemptive behavior ensures that no single goroutine can monopolize the CPU for an extended period.*
非協調型では、Goroutine が自ら制御を譲らなくても、Go ランタイムが強制的に一時停止させることができます。この仕組みによって、1つの Goroutine が延々と CPU を独占することが防がれます。

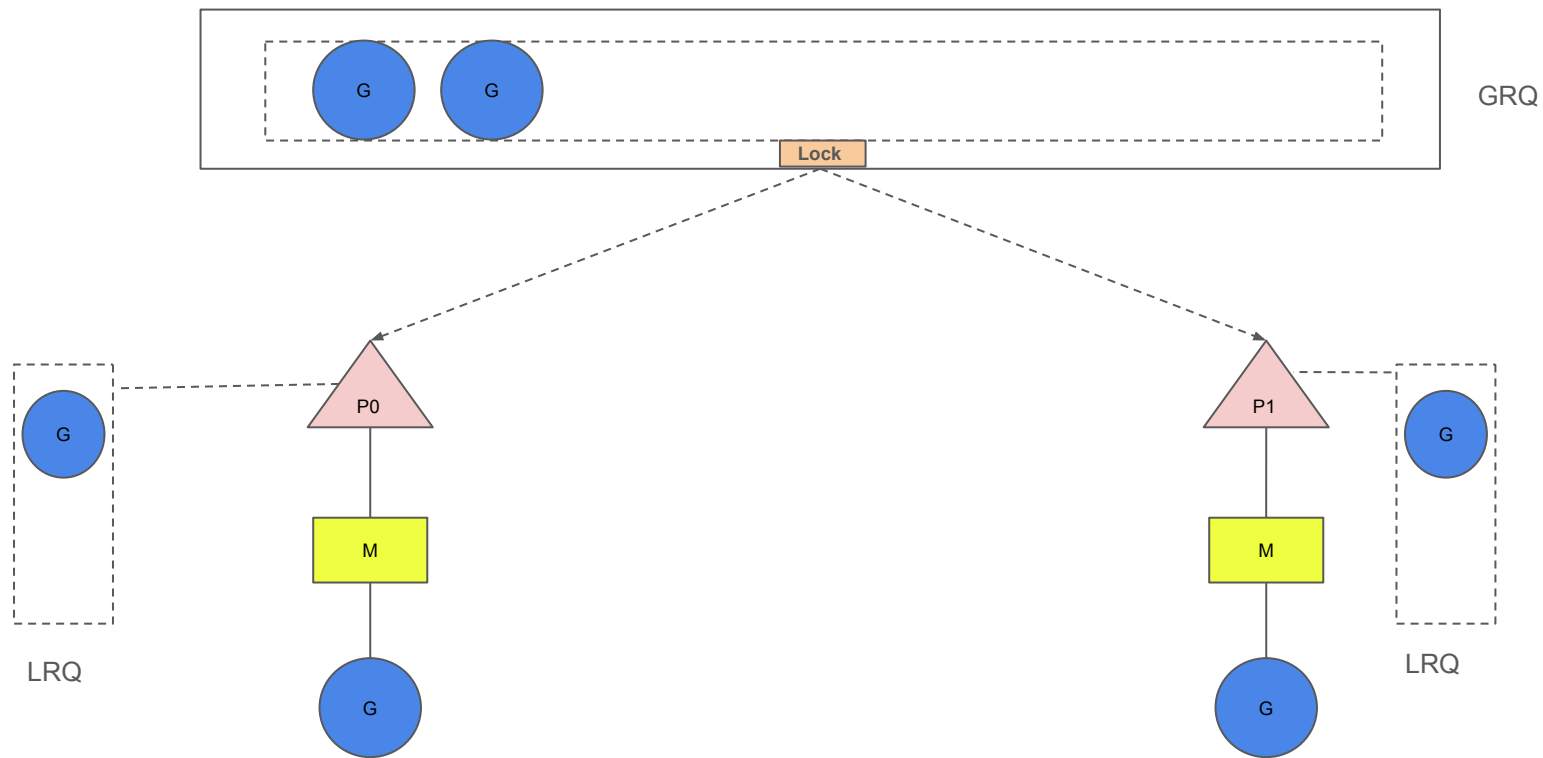


Been running for 10 ms

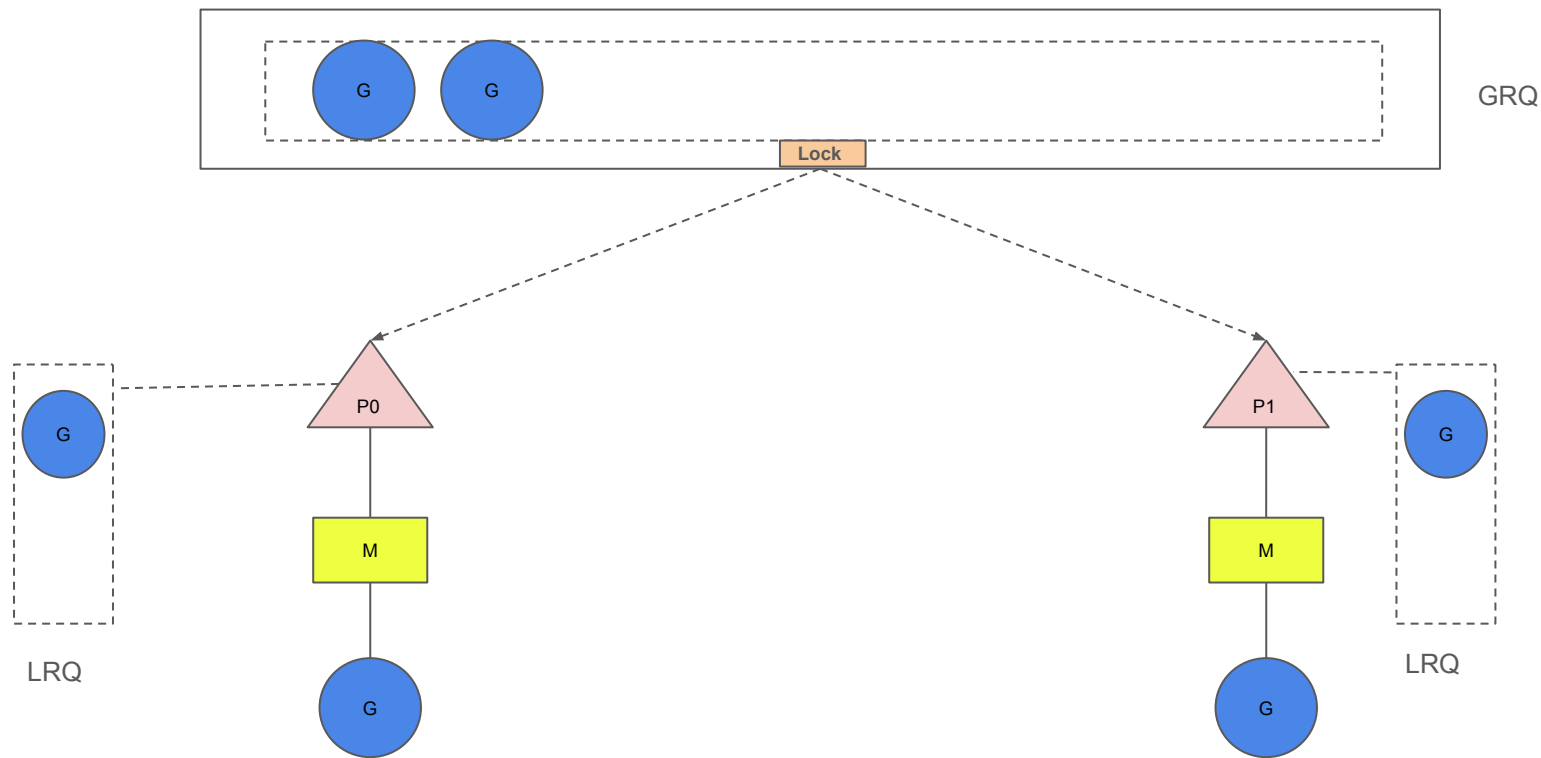
Sysmon Daemon



Where does the preempted go routine ends up going 🤔?

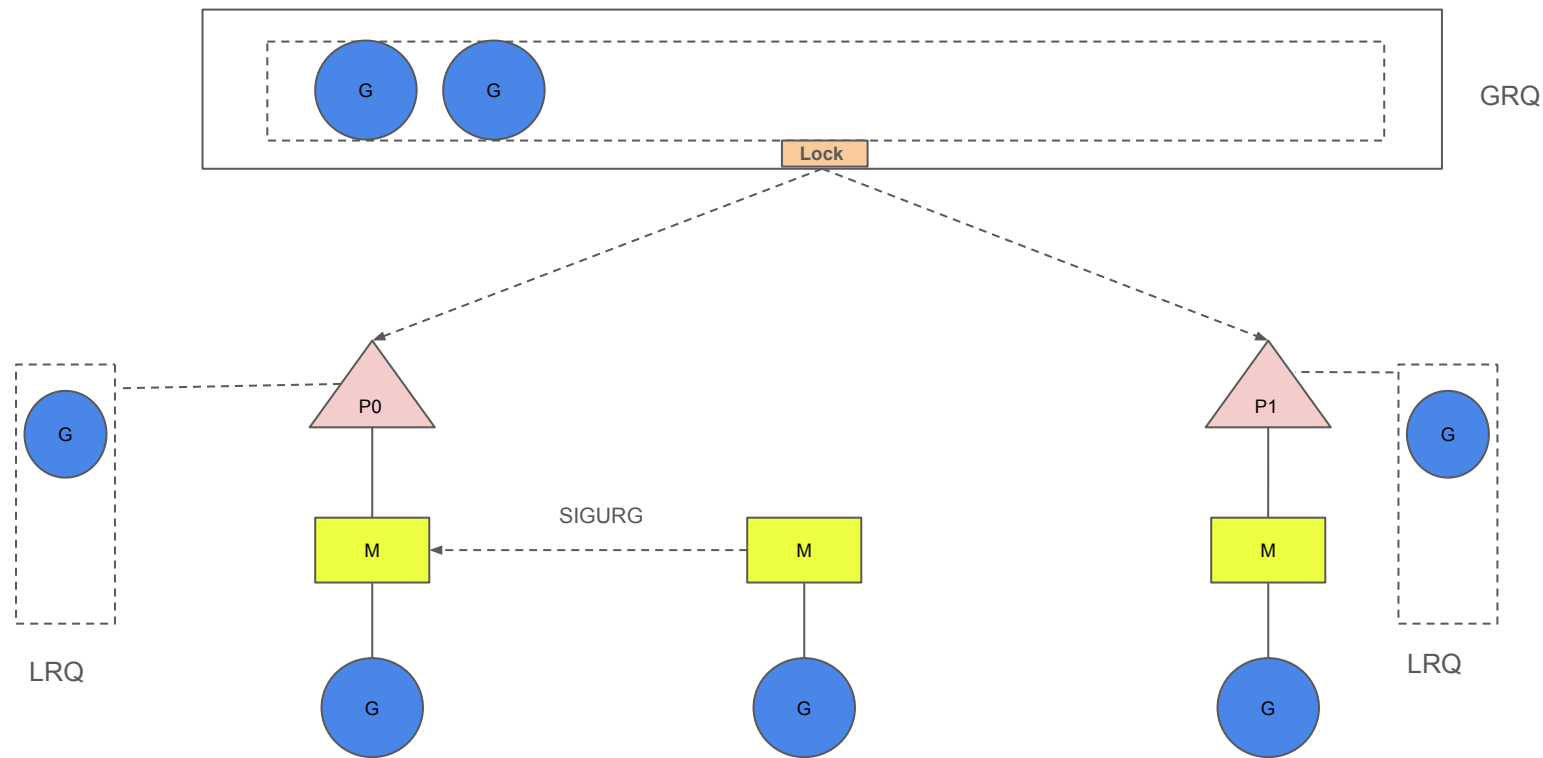


GRQ = Global Run Queue
LRQ = Local Run Queue



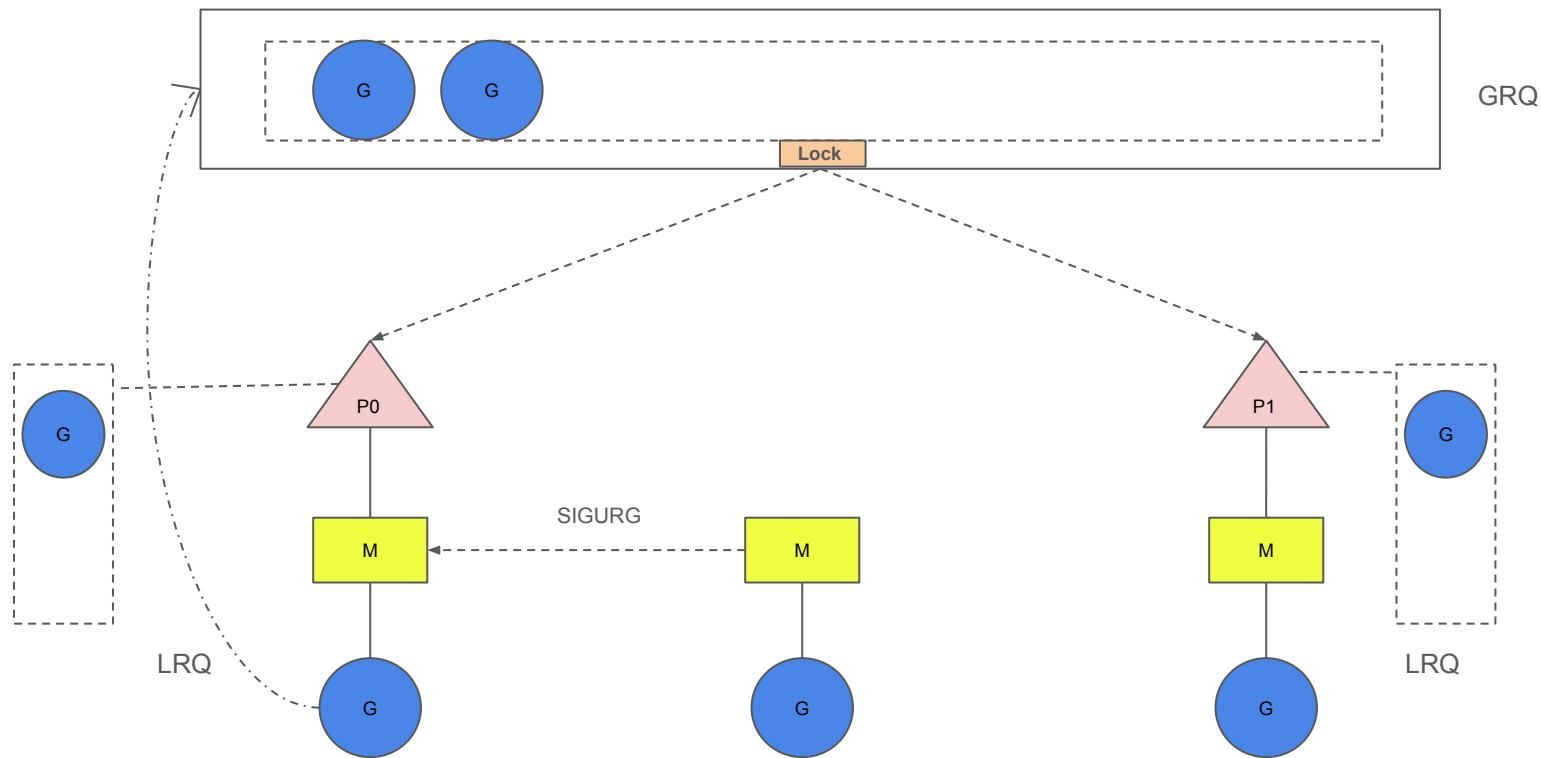
Long running goroutines

GRQ = Global Run Queue
LRQ = Local Run Queue



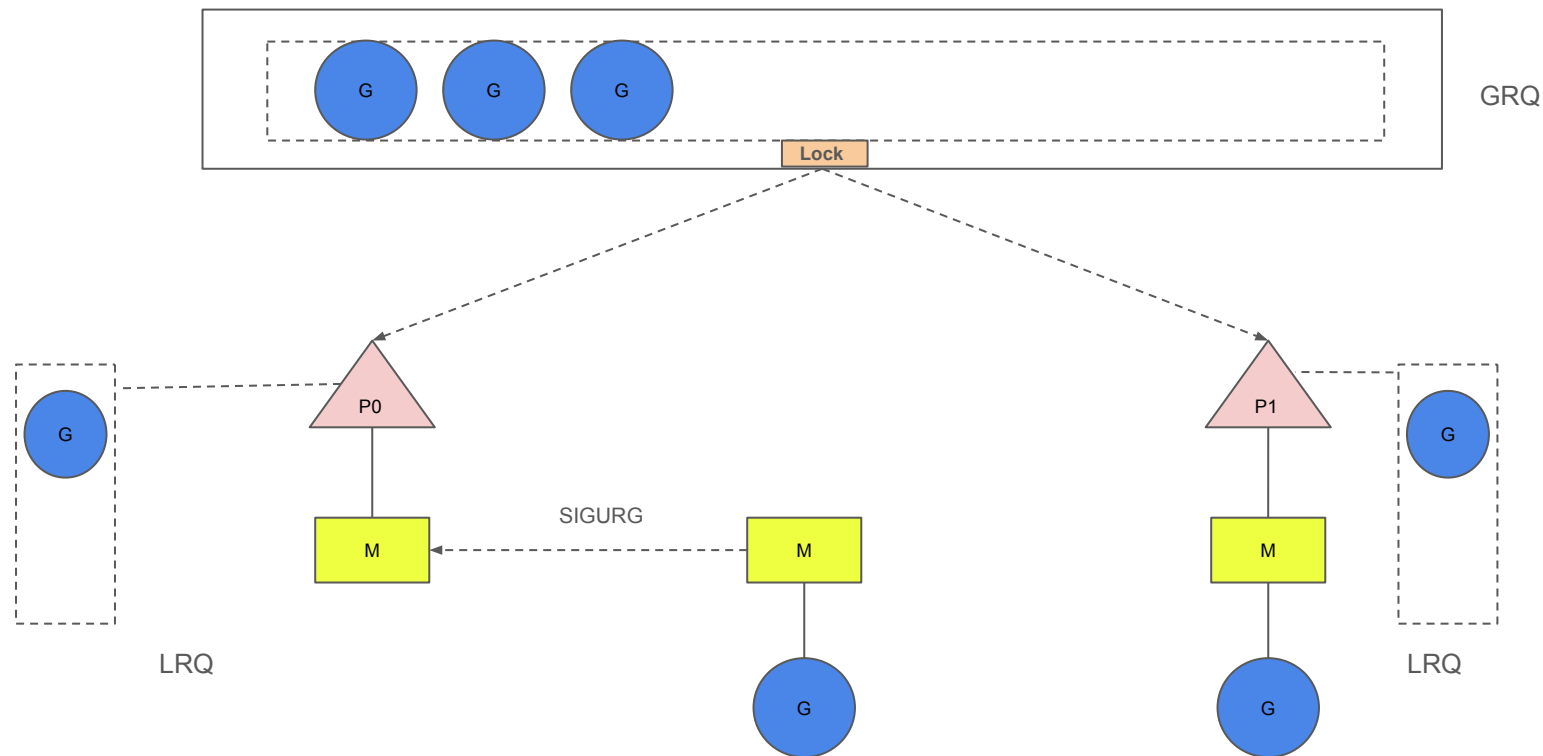
Long running goroutines

GRQ = Global Run Queue
LRQ = Local Run Queue

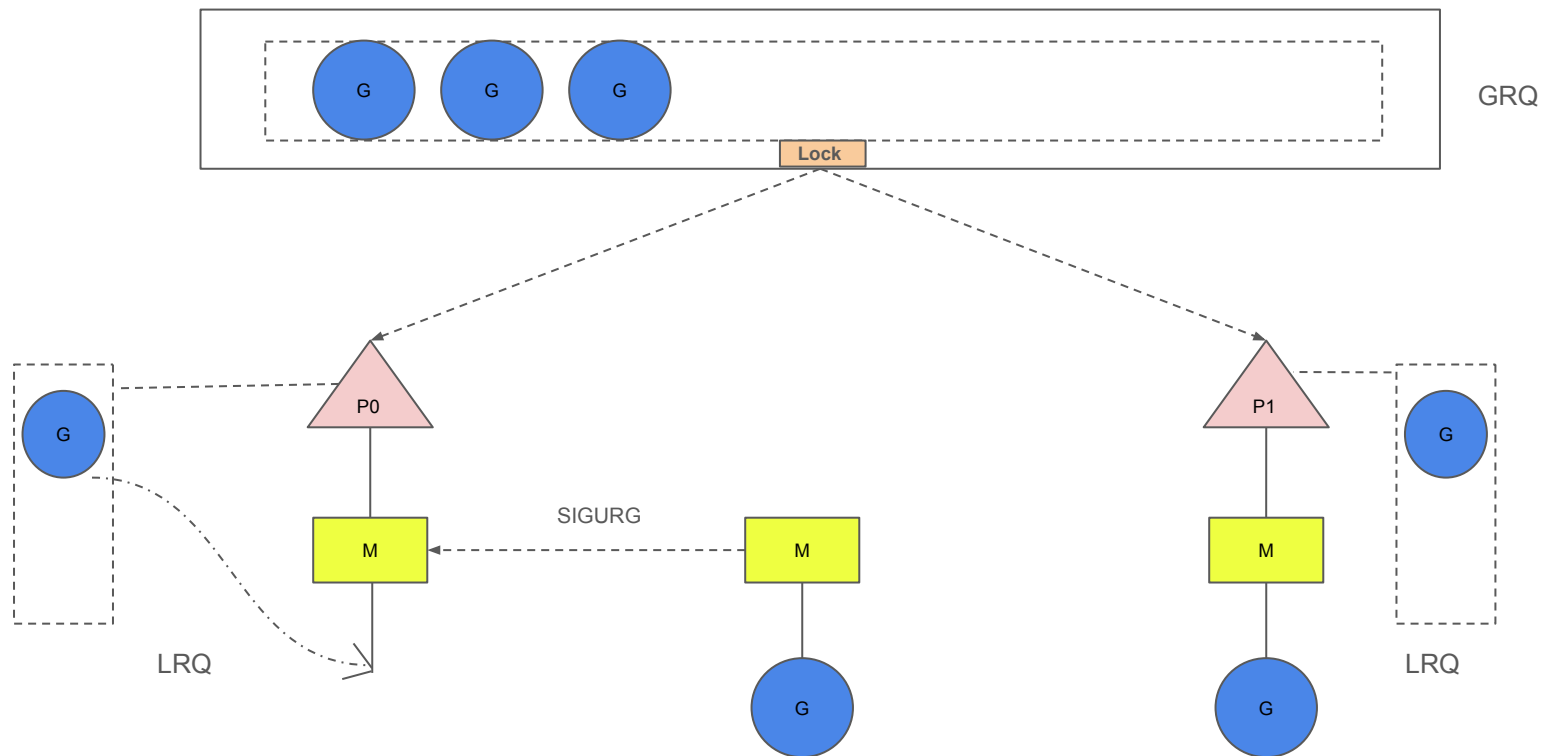


Long running goroutines

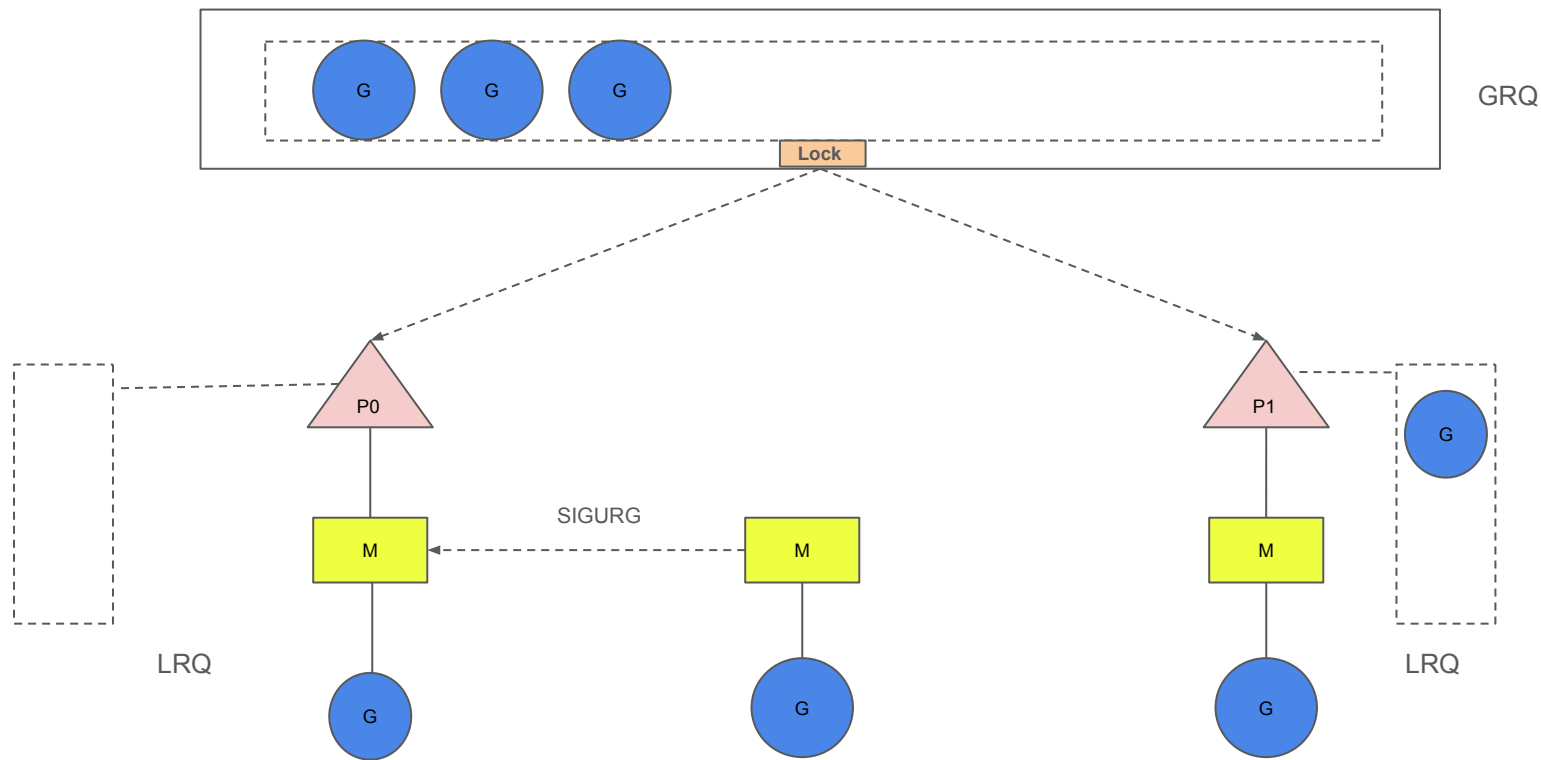
GRQ = Global Run Queue
LRQ = Local Run Queue



GRQ = Global Run Queue
LRQ = Local Run Queue




GRQ = Global Run Queue
LRQ = Local Run Queue




GRQ = Global Run Queue
LRQ = Local Run Queue

Visualization



Visualize scheduling, preemption, and runnable queues with
runtime/trace




```
func cpuBound(id int, iters int, wg *sync.WaitGroup) {  
    defer wg.Done()  
    sum := 0  
    for i := 0; i < iters; i++ {  
        sum += i ^ (i >> 3)  
        if i%500000 == 0 {  
            // Hint to runtime: allow preemption points  
            runtime.Gosched()  
        }  
    }  
    _ = sum  
}
```

- *Tight arithmetic to consume CPU*
- *The inner loop does simple integer ops to keep the compiler from optimizing away work*
- *runtime.Gosched() adds explicit yield points*

```
func ioBlocked(id int, wg *sync.WaitGroup) {  
→   defer wg.Done()  
→   // Simulate "blocking" on sleep to show goroutine parking/unparking  
→   time.Sleep(300 * time.Millisecond)  
}
```

- *ioBlocked Function Just sleeps for 300 ms to simulate a blocking operation*

```
func main() {
    // Encourage multiple Ps to see scheduling across threads
    runtime.GOMAXPROCS(4)

    f, err := os.Create("trace.out")
    if err != nil {
        panic(err)
    }
    defer f.Close()

    if err := trace.Start(f); err != nil {
        panic(err)
    }
    defer trace.Stop()

    var wg sync.WaitGroup
    rand.Seed(time.Now().UnixNano())


    // Mix cpu-bound and "IO"-blocked goroutines
    cpuG := 6
    ioG := 4


    wg.Add(cpuG + ioG)
    for i := 0; i < cpuG; i++ {
        it := 3_000_000 + rand.Intn(1_000_000)
        go cpuBound(i, it, &wg)
    }
    for i := 0; i < ioG; i++ {
        go ioBlocked(i, &wg)
    }

    wg.Wait()
    fmt.Println("done; trace written to trace.out")
}
```

- *The program mixes CPU-bound goroutines and blocking sleepers*
- *Then records a runtime execution trace so the scheduler's behavior, park/unpark, and preemption are visible in the trace UI.*


Demonstration





- The 4 Ps thanks to GOMAXPROCS(4) colored slices are goroutines running on a P's M, and blank gaps mean the P had nothing runnable or we're between events.
- The sleepers park on time. Sleep see them disappear from Ps then a timer wakes them and they become Runnable and run again, that's unpark.
- Thee cpu Bound goroutines yield at Gosched, so we see short slices and frequent context switches
*removing Gosched shows longer runs until the runtime preempts (10 ms).
- Notice the same GID switching P lanes that Work Stealing .
- Preemption occurs either cooperatively (runtime.Gosched()) or asynchronously Removing Gosched shows the runtime-driven preemption more clearly as longer uninterrupted slices that get cut by async preemption .

Beginner Pitfalls & My Learnings



Blocking isn't just I/O

- *Network/disk/syscalls block a G*
- *Long CPU loops starve others*
- *Big buffers hide backpressure*
- *Actions: timeouts, contexts, small critical sections*

```
// Bad: long work under lock
mu.Lock()
data = heavyCompute(load()) // blocks others
mu.Unlock()

// Better: shrink the critical section
x := load()
y := heavyCompute(x) // outside lock
mu.Lock()
data = y
mu.Unlock()
```

GOMAXPROCS: Measure, Don't Guess

- Controls parallel goroutine execution
- CPU-bound: \approx NumCPU
- I/O-bound: too high \rightarrow context switching
- Start at default; tune from traces

```
// Inspect and set intentionally  
n := runtime.GOMAXPROCS(0)  
_ = n // current setting  
runtime.GOMAXPROCS(runtime.NumCPU())
```


Preemption: Give the Scheduler Air

- *Tight CPU loops can hog a P*
- *Insert calls/checks; use contexts*
- *Break big tasks into steps*


```
func crunch(ctx context.Context) {  
    for i := 0; i < bigN; i++ {  
        if i%1_000 == 0 {  
            select {  
            case <-ctx.Done():  
                return  
            default:  
            }  
        }  
        step(i) // function call = preemption point  
    }  
}
```

My Quick Fix Checklist

- *Where can this block? (I/O, locks, channels)*
- *Is concurrency bounded ?*
- *Are tasks too chunky?*
- *Are locks too coarse?*
- *Do traces show runnable goroutines waiting?*




Conclusion




Why Understanding the Scheduler Matters

- *Predictability under load*
- *Fewer “mystery slowdowns”*
- *Better decisions: pooling, buffering, timeouts*
- *Faster debugging with data, not guesses*




After go func(), the Scheduler Conducts

- *Make it easy: bounded, balanced, measurable*
- *Learn the patterns; trust the traces*
- *Takeaway: Measure, don't guess*



References :

- <https://community.sap.com/t5/additional-blog-posts-by-sap/mastering-concurrency-unveiling-the-magic-of-go-s-scheduler/ba-p/13577437>
- <https://go.dev/src/runtime/proc.go>
- <https://github.com/golang/proposal/blob/master/design/24543-non-cooperative-preemption.md>
- https://www.cs.columbia.edu/~aho/cs6998/reports/12-12-11_DeshpandeSponslerWeiss_GO.pdf
- <https://medium.com/@hatronix/inside-the-go-scheduler-a-step-by-step-look-at-goroutine-management-1a8cbe9d5dbd>
- <https://medium.com/a-journey-with-go/go-work-stealing-in-go-scheduler-d439231be64d>
- https://docs.google.com/document/d/1TTj4T2JO42uD5ID9e89oa0sLKhJYD0Y_kqxDv3l3XMw/edit?tab=t.0#heading=h.mmq8lm48qfcw
- https://www.youtube.com/watch?v=S-MaTH8WpOM&ab_channel=Hypermode



聞いてくれて
ありがとうございます



[Session Code Repo](#)



[Session Slides](#)

Q/A

